
 Personal Open source Business Explore Pricing Blog Support

This repositorySearchSign inSign up

 **ronghanghu / rcnn**
forked from rbgirshick/rcnn

Watch1Star0Fork392

Code

Pull requests0

Pulse


Graphs

R-CNN: Regions with Convolutional Neural Network Features




46 commits8 branches2 releases2 contributors

Branch: masterNew pull requestNew fileFind fileHTTPShttps://github.com/ronghaDownload ZIP

This branch is 3 commits behind rbgirshick:master.Pull requestCompare

 **rbgirshick** add notes about these functions which are not useful to the general p... Latest commit 22f1605 on 15 Jul 2014

bbox_regression	load cached results if they exist	2 years ago
bin	add missing bin directory	2 years ago
cachedir	make cachedir setup same with a local override file	2 years ago
data	update data download READMEs and main README	2 years ago
datasets	initial checkin	2 years ago
examples	add threshold to demo	2 years ago
experiments	show how to use RCNN_CONFIG_OVERRIDE	2 years ago
external	cleanup installation instructions after a walkthrough; remove externa...	2 years ago
feat_cache	add feat_cache/README.md	2 years ago
finetuning	upgrade prototxt to caffe proto v1	2 years ago
imdb	nms thresh as argument	2 years ago
model-defs	improve demo; update models to caffe's v1 proto messages	2 years ago
nms	* fix caffe net paths	2 years ago
selective_search	add notes about these functions which are not useful to the general p...	2 years ago
utils	add one-off function for testing with the built-in softmax	2 years ago
vis	fix incorrect receptive field size	2 years ago
.gitignore	update data download READMEs and main README	2 years ago
LICENSE	Update LICENSE	2 years ago
README.md	fix git clone error	2 years ago
rcnn_build.m	initial checkin	2 years ago
rcnn_cache_pool5_features.m	sprinkle functions with basic documentation	2 years ago
rcnn_config.m	sprinkle functions with basic documentation	2 years ago
rcnn_config_local.example.m	make cachedir setup same with a local override file	2 years ago
rcnn_create_model.m	add copyright notices in files	2 years ago
rcnn_extract_regions.m	fix bug reported by Shubham Tulsiani	2 years ago
rcnn_feature_stats.m	add copyright notices in files	2 years ago
rcnn_features.m	allow for features from new and old matcaffe by canonicalizing them into	2 years ago
rcnn_im_crop.m	sprinkle functions with basic documentation	2 years ago
rcnn_load_cached_pool5_features.m	sprinkle functions with basic documentation	2 years ago
rcnn_load_model.m	sprinkle functions with basic documentation	2 years ago
rcnn_pool5_to_fcX.m	sprinkle functions with basic documentation	2 years ago
rcnn_scale_features.m	sprinkle functions with basic documentation	2 years ago

 rcnn_test.m	sprinkle functions with basic documentation	2 years ago
 rcnn_train.m	r2013a compatibility fix	2 years ago
 startup.m	improve demo; update models to caffe's v1 proto messages	2 years ago

README.md

R-CNN: *Regions with Convolutional Neural Network Features*

Created by Ross Girshick, Jeff Donahue, Trevor Darrell and Jitendra Malik at UC Berkeley EECS.

Acknowledgements: a huge thanks to Yangqing Jia for creating Caffe and the BVLC team, with a special shoutout to Evan Shelhamer, for maintaining Caffe and helping to merge the R-CNN fine-tuning code into Caffe.

Introduction

R-CNN is a state-of-the-art visual object detection system that combines bottom-up region proposals with rich features computed by a convolutional neural network. At the time of its release, R-CNN improved the previous best detection performance on PASCAL VOC 2012 by 30% relative, going from 40.9% to 53.3% mean average precision. Unlike the previous best results, R-CNN achieves this performance without using contextual rescoring or an ensemble of feature types.

R-CNN was initially described in an [arXiv tech report](#) and will appear in a forthcoming CVPR 2014 paper.

Citing R-CNN

If you find R-CNN useful in your research, please consider citing:

```
@inproceedings{girshick14CVPR,
  Author = {Girshick, Ross and Donahue, Jeff and Darrell, Trevor and Malik, Jitendra},
  Title = {Rich feature hierarchies for accurate object detection and semantic segmentation},
  Booktitle = {Computer Vision and Pattern Recognition},
  Year = {2014}
}
```

License

R-CNN is released under the Simplified BSD License (refer to the LICENSE file for details).

PASCAL VOC detection results

Method	VOC 2007 mAP	VOC 2010 mAP	VOC 2012 mAP
R-CNN	54.2%	50.2%	49.6%
R-CNN bbox reg	58.5%	53.7%	53.3%

- VOC 2007 per-class results are available in our [CVPR14 paper](#)
- VOC 2010 per-class results are available on the [VOC 2010 leaderboard](#)
- VOC 2012 per-class results are available on the [VOC 2012 leaderboard](#)
- These models are available in the model package (see below)

ImageNet 200-class detection results

Method	ILSVRC2013 test mAP
R-CNN bbox reg	31.4%

- For more details see the updated [R-CNN tech report](#) (Sections 2.5 and 4, in particular)
- This model is available in the model package (see below)
- The code that was used for training is in the `ilsvrc` branch (still needs some cleanup before merging into `master`)

Installing R-CNN

1. Prerequisites

- i. MATLAB (tested with 2012b on 64-bit Linux)
- ii. Caffe's [prerequisites](#)

2. Install Caffe (this is the most complicated part)

- i. R-CNN has been checked for compatability against Caffe release v0.999 (kona-snow), however it *should* also work with the current Caffe master
- ii. Download [Caffe v0.999](#)
- iii. Follow the [Caffe installation instructions](#)
- iv. Let's call the place where you installed caffe `$CAFFE_ROOT` (you can run `export CAFFE_ROOT=$(pwd)`)
- v. **Important:** Make sure to compile the Caffe MATLAB wrapper, which is not built by default: `make matcaffe`
- vi. **Important:** Make sure to run `cd $CAFFE_ROOT/data/ilsrvrc12 && ./get_ilsrvrc_aux.sh` to download the ImageNet image mean

3. Install R-CNN

- i. Get the R-CNN source code by cloning the repository: `git clone https://github.com/rbgirshick/rcnn.git`
- ii. Now change into the R-CNN source code directory: `cd rcnn`
- iii. R-CNN expects to find Caffe in `external/caffe` , so create a symlink: `ln -sf $CAFFE_ROOT external/caffe`
- iv. Start MATLAB (make sure you're still in the `rcnn` directory): `matlab`
- v. You'll be prompted to download the [Selective Search](#) code, which we cannot redistribute. Afterwards, you should see the message `R-CNN startup done` followed by the MATLAB prompt `>> .`
- vi. Run the build script: `>> rcnn_build()` (builds `liblinear` and [Selective Search](#)). Don't worry if you see compiler warnings while building `liblinear`, this is normal on my system.
- vii. Check that Caffe and MATLAB wrapper are set up correctly (this code should run without error): `>> key = caffe('get_init_key');` (expected output is `key = -2`)
- viii. Download the model package, which includes precompute models (see below).

Common issues: You may need to set an `LD_LIBRARY_PATH` before you start MATLAB. If you see a message like "Invalid MEX-file 'path/to/rcnn/external/caffe/matlab/caffe/caffe.mexa64': libmkl_rt.so: cannot open shared object file: No such file or directory" then make sure that CUDA and MKL are in your `LD_LIBRARY_PATH` . On my system, I use:

```
export LD_LIBRARY_PATH=/opt/intel/mkl/lib/intel64:/usr/local/cuda/lib64
```

Downloading pre-computed models (the model package)

The quickest way to get started is to download pre-computed R-CNN detectors. Currently we have detectors trained on PASCAL VOC 2007 train+val, 2012 train, and ILSVRC13 train+val. Unfortunately the download is large (1.5GB), so brew some coffee or take a walk while waiting.

From the `rcnn` folder, run the model fetch script: `./data/fetch_models.sh` .

This will populate the `rcnn/data` folder with `caffe_nets` and `rcnn_models` . See `rcnn/data/README.md` for details.

Pre-computed selective search boxes can also be downloaded for VOC2007, VOC2012, and ILSVRC13. From the `rcnn` folder, run the selective search data fetch script: `./data/fetch_selective_search_data.sh` .

This will populate the `rcnn/data` folder with `selective_selective_data` .

Caffe compatibility note: R-CNN has been updated to use the new Caffe proto messages that were rolled out in Caffe v0.999. The model package contains models in the up-to-date proto format. If, for some reason, you need to get the old (Caffe proto v0) models, they can still be downloaded: [VOC models](#) [ILSVRC13 model](#).

Running an R-CNN detector on an image

Let's assume that you've downloaded the precomputed detectors. Now:

1. Change to where you installed R-CNN: `cd rcnn` .
2. Start MATLAB `matlab` .
 - o **Important:** if you don't see the message `R-CNN startup done` when MATLAB starts, then you probably didn't start MATLAB in `rcnn` directory.
3. Run the demo: `>> rcnn_demo`

4. Enjoy the detected bicycle and person

Training your own R-CNN detector on PASCAL VOC

Let's use PASCAL VOC 2007 as an example. The basic pipeline is:

```
extract features to disk -> train SVMs -> test
```

You'll need about 200GB of disk space free for the feature cache (which is stored in `rcnn/feat_cache` by default; symlink `rcnn/feat_cache` elsewhere if needed). **It's best if the feature cache is on a fast, local disk.** Before running the pipeline, we first need to install the PASCAL VOC 2007 dataset.

Installing PASCAL VOC 2007

1. Download the training, validation, test data and VOCdevkit:

```
wget http://pascallin.ecs.soton.ac.uk/challenges/VOC/voc2007/VOCtrainval_06-Nov-2007.tar
wget http://pascallin.ecs.soton.ac.uk/challenges/VOC/voc2007/VOCtest_06-Nov-2007.tar
wget http://pascallin.ecs.soton.ac.uk/challenges/VOC/voc2007/VOCdevkit_08-Jun-2007.tar
```

2. Extract all of these tars into one directory, it's called `VOCdevkit`.

```
tar xvf VOCtrainval_06-Nov-2007.tar
tar xvf VOCtest_06-Nov-2007.tar
tar xvf VOCdevkit_08-Jun-2007.tar
```

3. It should have this basic structure:

```
VOCdevkit/                                % development kit
VOCdevkit/VOCcode/                        % VOC utility code
VOCdevkit/VOC2007                         % image sets, annotations, etc.
... and several other directories ...
```

4. I use a symlink to hook the R-CNN codebase to the PASCAL VOC dataset:

```
ln -sf /your/path/to/voc2007/VOCdevkit /path/to/rcnn/datasets/VOCdevkit2007
```

Extracting features

```
>> rcnn_exp_cache_features('train'); % chunk1
>> rcnn_exp_cache_features('val'); % chunk2
>> rcnn_exp_cache_features('test_1'); % chunk3
>> rcnn_exp_cache_features('test_2'); % chunk4
```

Pro tip: on a machine with one hefty GPU (e.g., k20, k40, titan) and a six-core processor, I run start two MATLAB sessions each with a three worker matlabpool. I then run chunk1 and chunk2 in parallel on that machine. In this setup, completing chunk1 and chunk2 takes about 8-9 hours (depending on your CPU/GPU combo and disk) on a single machine. Obviously, if you have more machines you can hack this function to split the workload.

Training R-CNN models and testing

Now to run the training and testing code, use the following experiments script:

```
>> test_results = rcnn_exp_train_and_test()
```

Note: The training and testing procedures save models and results under `rcnn/cachedir` by default. You can customize this by creating a local config file named `rcnn_config_local.m` and defining the experiment directory variable `EXP_DIR`. Look at `rcnn_config_local.example.m` for an example.

Training an R-CNN detector on another dataset

It should be easy to train an R-CNN detector using another detection dataset as long as that dataset has *complete* bounding box annotations (i.e., all instances of all classes are labeled).

To support a new dataset, you define three functions: (1) one that returns a structure that describes the class labels and list of images; (2) one that returns a region of interest (roi) structure that describes the bounding box annotations; and (3) one that provides an test evaluation function.

You can follow the PASCAL VOC implementation as your guide:

- `imdb/imdb_from_voc.m` (list of images and classes)
- `imdb/roidb_from_voc.m` (region of interest database)
- `imdb/imdb_eval_voc.m` (evaluation)

Fine-tuning a CNN for detection with Caffe

As an example, let's see how you would fine-tune a CNN for detection on PASCAL VOC 2012.

1. Create window files for VOC 2012 train and VOC 2012 val.
 - i. Start MATLAB in the `rcnn` directory
 - ii. Get the imdb for VOC 2012 train: `>> imdb_train = imdb_from_voc('datasets/VOCdevkit2012', 'train', '2012');`
 - iii. Get the imdb for VOC 2012 val: `>> imdb_val = imdb_from_voc('datasets/VOCdevkit2012', 'val', '2012');`
 - iv. Create the window file for VOC 2012 train: `>> rcnn_make_window_file(imdb_train, 'external/caffe/examples/pascal-finetuning');`
 - v. Create the window file for VOC 2012 val: `>> rcnn_make_window_file(imdb_val, 'external/caffe/examples/pascal-finetuning');`
 - vi. Exit MATLAB
2. Run fine-tuning with Caffe
 - i. Copy the fine-tuning prototxt files: `cp finetuning/voc_2012_prototxt/pascal_finetune_* external/caffe/examples/pascal-finetuning/`
 - ii. Change directories to `external/caffe/examples/pascal-finetuning`
 - iii. Execute the fine-tuning code (make sure to replace `/path/to/rcnn` with the actual path to where R-CNN is installed):

```
GL0G_logtostderr=1 ../../build/tools/finetune_net.bin \
pascal_finetune_solver.prototxt \
/path/to/rcnn/data/caffe_nets/ilsrvrc_2012_train_iter_310k 2>&1 | tee log.txt
```

Note: In my experiments, I've let fine-tuning run for 70k iterations, although with hindsight it appears that improvement in mAP saturates at around 40k iterations.