

caffe (/github/BVLC/caffe/tree/master) / examples (/github/BVLC/caffe/tree/master/examples)

## Brewing Logistic Regression then Going Deeper

While Caffe is made for deep networks it can likewise represent "shallow" models like logistic regression for classification. We'll do simple logistic regression on synthetic data that we'll generate and save to HDF5 to feed vectors to Caffe. Once that model is done, we'll add layers to improve accuracy. That's what Caffe is about: define a model, experiment, and then deploy.

In [1]:

```
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline

import os
os.chdir('.')

import sys
sys.path.insert(0, './python')
import caffe

import os
import h5py
import shutil
import tempfile

import sklearn
import sklearn.datasets
import sklearn.linear_model

import pandas as pd
```


Synthesize a dataset of 10,000 4-vectors for binary classification with 2 informative features and 2 noise features.

In [2]:

```
X, y = sklearn.datasets.make_classification(
    n_samples=10000, n_features=4, n_redundant=0, n_informative=2,
    n_clusters_per_class=2, hypercube=False, random_state=0
)

# Split into train and test
X, Xt, y, yt = sklearn.cross_validation.train_test_split(X, y)

# Visualize sample of the data
ind = np.random.permutation(X.shape[0])[:1000]
df = pd.DataFrame(X[ind])
_ = pd.scatter_matrix(df, figsize=(9, 9), diagonal='kde', marker='o', s=40, alpha=.4, c=
y[ind])
```



Learn and evaluate scikit-learn's logistic regression with stochastic gradient descent (SGD) training. Time and check the classifier's accuracy.

In [3]:

```
%%timeit
# Train and test the scikit-learn SGD logistic regression.
clf = sklearn.linear_model.SGDClassifier(
    loss='log', n_iter=1000, penalty='l2', alpha=5e-4, class_weight='auto')

clf.fit(X, y)
yt_pred = clf.predict(Xt)
print('Accuracy: {:.3f}'.format(sklearn.metrics.accuracy_score(yt, yt_pred)))
```

```
Accuracy: 0.781
Accuracy: 0.781
Accuracy: 0.781
Accuracy: 0.781
1 loop, best of 3: 372 ms per loop
```

Save the dataset to HDF5 for loading in Caffe.

In [4]:

```
# Write out the data to HDF5 files in a temp directory.
# This file is assumed to be caffe_root/examples/hdf5_classification.ipynb
dirname = os.path.abspath('./examples/hdf5_classification/data')
if not os.path.exists(dirname):
    os.makedirs(dirname)

train_filename = os.path.join(dirname, 'train.h5')
test_filename = os.path.join(dirname, 'test.h5')

# HDF5DataLayer source should be a file containing a list of HDF5 filenames.
# To show this off, we'll list the same data file twice.
with h5py.File(train_filename, 'w') as f:
    f['data'] = X
    f['label'] = y.astype(np.float32)
with open(os.path.join(dirname, 'train.txt'), 'w') as f:
    f.write(train_filename + '\n')
    f.write(train_filename + '\n')

# HDF5 is pretty efficient, but can be further compressed.
comp_kwargs = {'compression': 'gzip', 'compression_opts': 1}
with h5py.File(test_filename, 'w') as f:
    f.create_dataset('data', data=Xt, **comp_kwargs)
    f.create_dataset('label', data=yt.astype(np.float32), **comp_kwargs)
with open(os.path.join(dirname, 'test.txt'), 'w') as f:
    f.write(test_filename + '\n')
```

Let's define logistic regression in Caffe through Python net specification. This is a quick and natural way to define nets that sidesteps manually editing the protobuf model.

In [5]:

```
from caffe import layers as L
from caffe import params as P

def logreg(hdf5, batch_size):
    # logistic regression: data, matrix multiplication, and 2-class softmax loss
    n = caffe.NetSpec()
    n.data, n.label = L.HDF5Data(batch_size=batch_size, source=hdf5, ntop=2)
    n.ip1 = L.InnerProduct(n.data, num_output=2, weight_filler=dict(type='xavier'))
    n.accuracy = L.Accuracy(n.ip1, n.label)
    n.loss = L.SoftmaxWithLoss(n.ip1, n.label)
    return n.to_proto()

train_net_path = 'examples/hdf5_classification/logreg_auto_train.prototxt'
with open(train_net_path, 'w') as f:
    f.write(str(logreg('examples/hdf5_classification/data/train.txt', 10)))

test_net_path = 'examples/hdf5_classification/logreg_auto_test.prototxt'
with open(test_net_path, 'w') as f:
    f.write(str(logreg('examples/hdf5_classification/data/test.txt', 10)))
```

Now, we'll define our "solver" which trains the network by specifying the locations of the train and test nets we defined above, as well as setting values for various parameters used for learning, display, and "snapshotting".

In [6]:

```

from caffe.proto import caffe_pb2

def solver(train_net_path, test_net_path):
    s = caffe_pb2.SolverParameter()

    # Specify locations of the train and test networks.
    s.train_net = train_net_path
    s.test_net.append(test_net_path)

    s.test_interval = 1000 # Test after every 1000 training iterations.
    s.test_iter.append(250) # Test 250 "batches" each time we test.

    s.max_iter = 10000      # # of times to update the net (training iterations)

    # Set the initial learning rate for stochastic gradient descent (SGD).
    s.base_lr = 0.01

    # Set `lr_policy` to define how the learning rate changes during training.
    # Here, we 'step' the learning rate by multiplying it by a factor `gamma`
    # every `stepsize` iterations.
    s.lr_policy = 'step'
    s.gamma = 0.1
    s.stepsize = 5000

    # Set other optimization parameters. Setting a non-zero `momentum` takes a
    # weighted average of the current gradient and previous gradients to make
    # learning more stable. L2 weight decay regularizes learning, to help prevent
    # the model from overfitting.
    s.momentum = 0.9
    s.weight_decay = 5e-4

    # Display the current training loss and accuracy every 1000 iterations.
    s.display = 1000

    # Snapshots are files used to store networks we've trained. Here, we'll
    # snapshot every 10K iterations -- just once at the end of training.
    # For larger networks that take longer to train, you may want to set
    # snapshot < max_iter to save the network and training state to disk during
    # optimization, preventing disaster in case of machine crashes, etc.
    s.snapshot = 10000
    s.snapshot_prefix = 'examples/hdf5_classification/data/train'

    # We'll train on the CPU for fair benchmarking against scikit-learn.
    # Changing to GPU should result in much faster training!
    s.solver_mode = caffe_pb2.SolverParameter.CPU

    return s

solver_path = 'examples/hdf5_classification/logreg_solver.prototxt'
with open(solver_path, 'w') as f:
    f.write(str(solver(train_net_path, test_net_path)))

```

Time to learn and evaluate our Caffeinated logistic regression in Python.

In [7]:

```
%%timeit
caffe.set_mode_cpu()
solver = caffe.get_solver(solver_path)
solver.solve()

accuracy = 0
batch_size = solver.test_nets[0].blobs['data'].num
test_iters = int(len(Xt) / batch_size)
for i in range(test_iters):
    solver.test_nets[0].forward()
    accuracy += solver.test_nets[0].blobs['accuracy'].data
accuracy /= test_iters

print("Accuracy: {:.3f}".format(accuracy))
```

Accuracy: 0.770

Accuracy: 0.770

Accuracy: 0.770

Accuracy: 0.770

1 loop, best of 3: 195 ms per loop

Do the same through the command line interface for detailed output on the model and solving.

In [8]:

```
!./build/tools/caffe train -solver examples/hdf5_classification/logreg_solver.prototxt
```

```
I0224 00:32:03.232779    655 caffe.cpp:178] Use CPU.
I0224 00:32:03.391911    655 solver.cpp:48] Initializing solver from parameters:
train_net: "examples/hdf5_classification/logreg_auto_train.prototxt"
test_net: "examples/hdf5_classification/logreg_auto_test.prototxt"
test_iter: 250
test_interval: 1000
base_lr: 0.01
display: 1000
max_iter: 10000
lr_policy: "step"
gamma: 0.1
momentum: 0.9
weight_decay: 0.0005
stepsize: 5000
snapshot: 10000
snapshot_prefix: "examples/hdf5_classification/data/train"
solver_mode: CPU
I0224 00:32:03.392065    655 solver.cpp:81] Creating training net from train
_net file: examples/hdf5_classification/logreg_auto_train.prototxt
I0224 00:32:03.392215    655 net.cpp:49] Initializing net from parameters:
state {
  phase: TRAIN
}
layer {
  name: "data"
  type: "HDF5Data"
  top: "data"
  top: "label"
  hdf5_data_param {
    source: "examples/hdf5_classification/data/train.txt"
    batch_size: 10
  }
}
layer {
  name: "ip1"
  type: "InnerProduct"
  bottom: "data"
  top: "ip1"
  inner_product_param {
    num_output: 2
    weight_filler {
      type: "xavier"
    }
  }
}
layer {
  name: "accuracy"
  type: "Accuracy"
  bottom: "ip1"
  bottom: "label"
  top: "accuracy"
}
```

```

layer {
  name: "loss"
  type: "SoftmaxWithLoss"
  bottom: "ipl"
  bottom: "label"
  top: "loss"
}

```

```

I0224 00:32:03.392365 655 layer_factory.hpp:77] Creating layer data
I0224 00:32:03.392382 655 net.cpp:106] Creating Layer data
I0224 00:32:03.392395 655 net.cpp:411] data -> data
I0224 00:32:03.392423 655 net.cpp:411] data -> label
I0224 00:32:03.392442 655 hdf5_data_layer.cpp:79] Loading list of HDF5 fi
lenames from: examples/hdf5_classification/data/train.txt
I0224 00:32:03.392473 655 hdf5_data_layer.cpp:93] Number of HDF5 files: 2
I0224 00:32:03.393473 655 hdf5.cpp:32] Datatype class: H5T_FLOAT
I0224 00:32:03.393862 655 net.cpp:150] Setting up data
I0224 00:32:03.393884 655 net.cpp:157] Top shape: 10 4 (40)
I0224 00:32:03.393894 655 net.cpp:157] Top shape: 10 (10)
I0224 00:32:03.393901 655 net.cpp:165] Memory required for data: 200
I0224 00:32:03.393911 655 layer_factory.hpp:77] Creating layer label_data
_l_split
I0224 00:32:03.393924 655 net.cpp:106] Creating Layer label_data_l_split
I0224 00:32:03.393934 655 net.cpp:454] label_data_l_split <- label
I0224 00:32:03.393945 655 net.cpp:411] label_data_l_split -> label_data_l
_split_0
I0224 00:32:03.393956 655 net.cpp:411] label_data_l_split -> label_data_l
_split_1
I0224 00:32:03.393970 655 net.cpp:150] Setting up label_data_l_split
I0224 00:32:03.393978 655 net.cpp:157] Top shape: 10 (10)
I0224 00:32:03.393986 655 net.cpp:157] Top shape: 10 (10)
I0224 00:32:03.393995 655 net.cpp:165] Memory required for data: 280
I0224 00:32:03.394001 655 layer_factory.hpp:77] Creating layer ipl
I0224 00:32:03.394012 655 net.cpp:106] Creating Layer ipl
I0224 00:32:03.394021 655 net.cpp:454] ipl <- data
I0224 00:32:03.394029 655 net.cpp:411] ipl -> ipl
I0224 00:32:03.394311 655 net.cpp:150] Setting up ipl
I0224 00:32:03.394323 655 net.cpp:157] Top shape: 10 2 (20)
I0224 00:32:03.394331 655 net.cpp:165] Memory required for data: 360
I0224 00:32:03.394348 655 layer_factory.hpp:77] Creating layer ipl_ipl_0_
split
I0224 00:32:03.394358 655 net.cpp:106] Creating Layer ipl_ipl_0_split
I0224 00:32:03.394366 655 net.cpp:454] ipl_ipl_0_split <- ipl
I0224 00:32:03.394374 655 net.cpp:411] ipl_ipl_0_split -> ipl_ipl_0_split
_0
I0224 00:32:03.394386 655 net.cpp:411] ipl_ipl_0_split -> ipl_ipl_0_split
_1
I0224 00:32:03.394395 655 net.cpp:150] Setting up ipl_ipl_0_split
I0224 00:32:03.394404 655 net.cpp:157] Top shape: 10 2 (20)
I0224 00:32:03.394424 655 net.cpp:157] Top shape: 10 2 (20)
I0224 00:32:03.394443 655 net.cpp:165] Memory required for data: 520
I0224 00:32:03.394450 655 layer_factory.hpp:77] Creating layer accuracy
I0224 00:32:03.394462 655 net.cpp:106] Creating Layer accuracy
I0224 00:32:03.394479 655 net.cpp:454] accuracy <- ipl_ipl_0_split_0
I0224 00:32:03.394489 655 net.cpp:454] accuracy <- label_data_l_split_0
I0224 00:32:03.394497 655 net.cpp:411] accuracy -> accuracy

I0224 00:32:03.394510 655 net.cpp:150] Setting up accuracy
I0224 00:32:03.394536 655 net.cpp:157] Top shape: (1)

```



```

I0224 00:32:03.394543 655 net.cpp:165] Memory required for data: 524
I0224 00:32:03.394551 655 layer_factory.hpp:77] Creating layer loss
I0224 00:32:03.394562 655 net.cpp:106] Creating Layer loss
I0224 00:32:03.394569 655 net.cpp:454] loss <- ip1_ip1_0_split_1
I0224 00:32:03.394577 655 net.cpp:454] loss <- label_data_1_split_1
I0224 00:32:03.394587 655 net.cpp:411] loss -> loss
I0224 00:32:03.394603 655 layer_factory.hpp:77] Creating layer loss
I0224 00:32:03.394624 655 net.cpp:150] Setting up loss
I0224 00:32:03.394634 655 net.cpp:157] Top shape: (1)
I0224 00:32:03.394641 655 net.cpp:160] with loss weight 1
I0224 00:32:03.394659 655 net.cpp:165] Memory required for data: 528
I0224 00:32:03.394665 655 net.cpp:226] loss needs backward computation.
I0224 00:32:03.394673 655 net.cpp:228] accuracy does not need backward co
mputation.
I0224 00:32:03.394682 655 net.cpp:226] ip1_ip1_0_split needs backward com
putation.
I0224 00:32:03.394690 655 net.cpp:226] ip1 needs backward computation.
I0224 00:32:03.394697 655 net.cpp:228] label_data_1_split does not need b
ackward computation.
I0224 00:32:03.394706 655 net.cpp:228] data does not need backward comput
ation.
I0224 00:32:03.394712 655 net.cpp:270] This network produces output accur
acy
I0224 00:32:03.394721 655 net.cpp:270] This network produces output loss
I0224 00:32:03.394731 655 net.cpp:283] Network initialization done.
I0224 00:32:03.394804 655 solver.cpp:181] Creating test net (#0) specifie
d by test_net file: examples/hdf5_classification/logreg_auto_test.prototxt
I0224 00:32:03.394836 655 net.cpp:49] Initializing net from parameters:
state {
  phase: TEST
}
layer {
  name: "data"
  type: "HDF5Data"
  top: "data"
  top: "label"
  hdf5_data_param {
    source: "examples/hdf5_classification/data/test.txt"
    batch_size: 10
  }
}
layer {
  name: "ip1"
  type: "InnerProduct"
  bottom: "data"
  top: "ip1"
  inner_product_param {
    num_output: 2
    weight_filler {
      type: "xavier"
    }
  }
}
layer {
  name: "accuracy"

  type: "Accuracy"
  bottom: "ip1"

```

```

    bottom: "label"
    top: "accuracy"
  }
  layer {
    name: "loss"
    type: "SoftmaxWithLoss"
    bottom: "ipl"
    bottom: "label"
    top: "loss"
  }

```

```

I0224 00:32:03.394953 655 layer_factory.hpp:77] Creating layer data
I0224 00:32:03.394964 655 net.cpp:106] Creating Layer data
I0224 00:32:03.394973 655 net.cpp:411] data -> data
I0224 00:32:03.394984 655 net.cpp:411] data -> label
I0224 00:32:03.394994 655 hdf5_data_layer.cpp:79] Loading list of HDF5 fi
lenames from: examples/hdf5_classification/data/test.txt
I0224 00:32:03.395009 655 hdf5_data_layer.cpp:93] Number of HDF5 files: 1
I0224 00:32:03.395937 655 net.cpp:150] Setting up data
I0224 00:32:03.395953 655 net.cpp:157] Top shape: 10 4 (40)
I0224 00:32:03.395963 655 net.cpp:157] Top shape: 10 (10)
I0224 00:32:03.395970 655 net.cpp:165] Memory required for data: 200
I0224 00:32:03.395978 655 layer_factory.hpp:77] Creating layer label_data
_1_split
I0224 00:32:03.395989 655 net.cpp:106] Creating Layer label_data_1_split
I0224 00:32:03.395997 655 net.cpp:454] label_data_1_split <- label
I0224 00:32:03.396005 655 net.cpp:411] label_data_1_split -> label_data_1
_split_0
I0224 00:32:03.396016 655 net.cpp:411] label_data_1_split -> label_data_1
_split_1
I0224 00:32:03.396028 655 net.cpp:150] Setting up label_data_1_split
I0224 00:32:03.396036 655 net.cpp:157] Top shape: 10 (10)
I0224 00:32:03.396044 655 net.cpp:157] Top shape: 10 (10)
I0224 00:32:03.396051 655 net.cpp:165] Memory required for data: 280
I0224 00:32:03.396059 655 layer_factory.hpp:77] Creating layer ipl
I0224 00:32:03.396069 655 net.cpp:106] Creating Layer ipl
I0224 00:32:03.396075 655 net.cpp:454] ipl <- data
I0224 00:32:03.396085 655 net.cpp:411] ipl -> ipl
I0224 00:32:03.396100 655 net.cpp:150] Setting up ipl
I0224 00:32:03.396109 655 net.cpp:157] Top shape: 10 2 (20)
I0224 00:32:03.396116 655 net.cpp:165] Memory required for data: 360
I0224 00:32:03.396138 655 layer_factory.hpp:77] Creating layer ipl_ipl_0_
split
I0224 00:32:03.396148 655 net.cpp:106] Creating Layer ipl_ipl_0_split
I0224 00:32:03.396157 655 net.cpp:454] ipl_ipl_0_split <- ipl
I0224 00:32:03.396164 655 net.cpp:411] ipl_ipl_0_split -> ipl_ipl_0_split
_0
I0224 00:32:03.396174 655 net.cpp:411] ipl_ipl_0_split -> ipl_ipl_0_split
_1
I0224 00:32:03.396185 655 net.cpp:150] Setting up ipl_ipl_0_split
I0224 00:32:03.396194 655 net.cpp:157] Top shape: 10 2 (20)
I0224 00:32:03.396203 655 net.cpp:157] Top shape: 10 2 (20)
I0224 00:32:03.396209 655 net.cpp:165] Memory required for data: 520
I0224 00:32:03.396216 655 layer_factory.hpp:77] Creating layer accuracy
I0224 00:32:03.396225 655 net.cpp:106] Creating Layer accuracy
I0224 00:32:03.396234 655 net.cpp:454] accuracy <- ipl_ipl_0_split_0
I0224 00:32:03.396241 655 net.cpp:454] accuracy <- label_data_1_split_0
I0224 00:32:03.396250 655 net.cpp:411] accuracy -> accuracy

```

```

I0224 00:32:03.396260 655 net.cpp:150] Setting up accuracy
I0224 00:32:03.396270 655 net.cpp:157] Top shape: (1)
I0224 00:32:03.396276 655 net.cpp:165] Memory required for data: 524
I0224 00:32:03.396283 655 layer_factory.hpp:77] Creating layer loss
I0224 00:32:03.396291 655 net.cpp:106] Creating Layer loss
I0224 00:32:03.396299 655 net.cpp:454] loss <- ip1_ip1_0_split_1
I0224 00:32:03.396307 655 net.cpp:454] loss <- label_data_1_split_1
I0224 00:32:03.396317 655 net.cpp:411] loss -> loss
I0224 00:32:03.396327 655 layer_factory.hpp:77] Creating layer loss
I0224 00:32:03.396339 655 net.cpp:150] Setting up loss
I0224 00:32:03.396349 655 net.cpp:157] Top shape: (1)
I0224 00:32:03.396356 655 net.cpp:160]         with loss weight 1
I0224 00:32:03.396365 655 net.cpp:165] Memory required for data: 528
I0224 00:32:03.396373 655 net.cpp:226] loss needs backward computation.
I0224 00:32:03.396381 655 net.cpp:228] accuracy does not need backward co
mputation.
I0224 00:32:03.396389 655 net.cpp:226] ip1_ip1_0_split needs backward com
putation.
I0224 00:32:03.396396 655 net.cpp:226] ip1 needs backward computation.
I0224 00:32:03.396404 655 net.cpp:228] label_data_1_split does not need b
ackward computation.
I0224 00:32:03.396412 655 net.cpp:228] data does not need backward comput
ation.
I0224 00:32:03.396420 655 net.cpp:270] This network produces output accur
acy
I0224 00:32:03.396427 655 net.cpp:270] This network produces output loss
I0224 00:32:03.396437 655 net.cpp:283] Network initialization done.
I0224 00:32:03.396455 655 solver.cpp:60] Solver scaffolding done.
I0224 00:32:03.396473 655 caffe.cpp:219] Starting Optimization
I0224 00:32:03.396482 655 solver.cpp:280] Solving
I0224 00:32:03.396489 655 solver.cpp:281] Learning Rate Policy: step
I0224 00:32:03.396499 655 solver.cpp:338] Iteration 0, Testing net (#0)
I0224 00:32:03.932615 655 solver.cpp:406]         Test net output #0: accurac
y = 0.4268
I0224 00:32:03.932656 655 solver.cpp:406]         Test net output #1: loss =
1.33093 (* 1 = 1.33093 loss)
I0224 00:32:03.932723 655 solver.cpp:229] Iteration 0, loss = 1.06081
I0224 00:32:03.932737 655 solver.cpp:245]         Train net output #0: accura
cy = 0.4
I0224 00:32:03.932749 655 solver.cpp:245]         Train net output #1: loss =
1.06081 (* 1 = 1.06081 loss)
I0224 00:32:03.932765 655 sgd_solver.cpp:106] Iteration 0, lr = 0.01
I0224 00:32:03.945551 655 solver.cpp:338] Iteration 1000, Testing net (#0
)
I0224 00:32:03.948048 655 solver.cpp:406]         Test net output #0: accurac
y = 0.694
I0224 00:32:03.948065 655 solver.cpp:406]         Test net output #1: loss =
0.60406 (* 1 = 0.60406 loss)
I0224 00:32:03.948091 655 solver.cpp:229] Iteration 1000, loss = 0.505853
I0224 00:32:03.948102 655 solver.cpp:245]         Train net output #0: accura
cy = 0.7
I0224 00:32:03.948113 655 solver.cpp:245]         Train net output #1: loss =
0.505853 (* 1 = 0.505853 loss)
I0224 00:32:03.948122 655 sgd_solver.cpp:106] Iteration 1000, lr = 0.01
I0224 00:32:03.960741 655 solver.cpp:338] Iteration 2000, Testing net (#0
)
I0224 00:32:03.963214 655 solver.cpp:406]         Test net output #0: accurac

```

```
y = 0.7372
I0224 00:32:03.963249 655 solver.cpp:406] Test net output #1: loss =
0.595267 (* 1 = 0.595267 loss)
I0224 00:32:03.963276 655 solver.cpp:229] Iteration 2000, loss = 0.549211
I0224 00:32:03.963289 655 solver.cpp:245] Train net output #0: accuracy = 0.7
I0224 00:32:03.963299 655 solver.cpp:245] Train net output #1: loss =
0.549211 (* 1 = 0.549211 loss)
I0224 00:32:03.963309 655 sgd_solver.cpp:106] Iteration 2000, lr = 0.01
I0224 00:32:03.975945 655 solver.cpp:338] Iteration 3000, Testing net (#0
)
I0224 00:32:03.978435 655 solver.cpp:406] Test net output #0: accuracy = 0.7732
I0224 00:32:03.978451 655 solver.cpp:406] Test net output #1: loss =
0.594998 (* 1 = 0.594998 loss)
I0224 00:32:03.978884 655 solver.cpp:229] Iteration 3000, loss = 0.66133
I0224 00:32:03.978911 655 solver.cpp:245] Train net output #0: accuracy = 0.8
I0224 00:32:03.978932 655 solver.cpp:245] Train net output #1: loss =
0.66133 (* 1 = 0.66133 loss)
I0224 00:32:03.978950 655 sgd_solver.cpp:106] Iteration 3000, lr = 0.01
I0224 00:32:03.992017 655 solver.cpp:338] Iteration 4000, Testing net (#0
)
I0224 00:32:03.994509 655 solver.cpp:406] Test net output #0: accuracy = 0.694
I0224 00:32:03.994525 655 solver.cpp:406] Test net output #1: loss =
0.60406 (* 1 = 0.60406 loss)
I0224 00:32:03.994551 655 solver.cpp:229] Iteration 4000, loss = 0.505853
I0224 00:32:03.994562 655 solver.cpp:245] Train net output #0: accuracy = 0.7
I0224 00:32:03.994573 655 solver.cpp:245] Train net output #1: loss =
0.505853 (* 1 = 0.505853 loss)
I0224 00:32:03.994583 655 sgd_solver.cpp:106] Iteration 4000, lr = 0.01
I0224 00:32:04.007200 655 solver.cpp:338] Iteration 5000, Testing net (#0
)
I0224 00:32:04.009686 655 solver.cpp:406] Test net output #0: accuracy = 0.7372
I0224 00:32:04.009702 655 solver.cpp:406] Test net output #1: loss =
0.595267 (* 1 = 0.595267 loss)
I0224 00:32:04.009727 655 solver.cpp:229] Iteration 5000, loss = 0.549211
I0224 00:32:04.009738 655 solver.cpp:245] Train net output #0: accuracy = 0.7
I0224 00:32:04.009749 655 solver.cpp:245] Train net output #1: loss =
0.549211 (* 1 = 0.549211 loss)
I0224 00:32:04.009758 655 sgd_solver.cpp:106] Iteration 5000, lr = 0.001
I0224 00:32:04.022734 655 solver.cpp:338] Iteration 6000, Testing net (#0
)
I0224 00:32:04.025177 655 solver.cpp:406] Test net output #0: accuracy = 0.7824
I0224 00:32:04.025193 655 solver.cpp:406] Test net output #1: loss =
0.593367 (* 1 = 0.593367 loss)
I0224 00:32:04.025545 655 solver.cpp:229] Iteration 6000, loss = 0.654873
I0224 00:32:04.025562 655 solver.cpp:245] Train net output #0: accuracy = 0.7
I0224 00:32:04.025573 655 solver.cpp:245] Train net output #1: loss =
0.654873 (* 1 = 0.654873 loss)
I0224 00:32:04.025583 655 sgd_solver.cpp:106] Iteration 6000, lr = 0.001
```

```

I0224 00:32:04.038586 655 solver.cpp:338] Iteration 7000, Testing net (#0
)
I0224 00:32:04.041016 655 solver.cpp:406] Test net output #0: accurac
y = 0.7704
I0224 00:32:04.041033 655 solver.cpp:406] Test net output #1: loss =
0.593842 (* 1 = 0.593842 loss)
I0224 00:32:04.041059 655 solver.cpp:229] Iteration 7000, loss = 0.46611
I0224 00:32:04.041071 655 solver.cpp:245] Train net output #0: accura
cy = 0.6
I0224 00:32:04.041082 655 solver.cpp:245] Train net output #1: loss =
0.46611 (* 1 = 0.46611 loss)
I0224 00:32:04.041091 655 sgd_solver.cpp:106] Iteration 7000, lr = 0.001
I0224 00:32:04.053722 655 solver.cpp:338] Iteration 8000, Testing net (#0
)
I0224 00:32:04.056171 655 solver.cpp:406] Test net output #0: accurac
y = 0.7788
I0224 00:32:04.056187 655 solver.cpp:406] Test net output #1: loss =
0.592847 (* 1 = 0.592847 loss)
I0224 00:32:04.056213 655 solver.cpp:229] Iteration 8000, loss = 0.615126
I0224 00:32:04.056224 655 solver.cpp:245] Train net output #0: accura
cy = 0.8
I0224 00:32:04.056236 655 solver.cpp:245] Train net output #1: loss =
0.615126 (* 1 = 0.615126 loss)
I0224 00:32:04.056244 655 sgd_solver.cpp:106] Iteration 8000, lr = 0.001
I0224 00:32:04.068853 655 solver.cpp:338] Iteration 9000, Testing net (#0
)
I0224 00:32:04.071291 655 solver.cpp:406] Test net output #0: accurac
y = 0.7808
I0224 00:32:04.071307 655 solver.cpp:406] Test net output #1: loss =
0.593293 (* 1 = 0.593293 loss)
I0224 00:32:04.071650 655 solver.cpp:229] Iteration 9000, loss = 0.654997
I0224 00:32:04.071666 655 solver.cpp:245] Train net output #0: accura
cy = 0.7
I0224 00:32:04.071677 655 solver.cpp:245] Train net output #1: loss =
0.654998 (* 1 = 0.654998 loss)
I0224 00:32:04.071687 655 sgd_solver.cpp:106] Iteration 9000, lr = 0.001
I0224 00:32:04.084717 655 solver.cpp:456] Snapshotting to binary proto fi
le examples/hdf5_classification/data/train_iter_10000.caffemodel
I0224 00:32:04.084885 655 sgd_solver.cpp:273] Snapshotting solver state t
o binary proto file examples/hdf5_classification/data/train_iter_10000.solv
erstate
I0224 00:32:04.084960 655 solver.cpp:318] Iteration 10000, loss = 0.46650
5
I0224 00:32:04.084977 655 solver.cpp:338] Iteration 10000, Testing net (#
0)
I0224 00:32:04.087514 655 solver.cpp:406] Test net output #0: accurac
y = 0.77
I0224 00:32:04.087532 655 solver.cpp:406] Test net output #1: loss =
0.593815 (* 1 = 0.593815 loss)
I0224 00:32:04.087541 655 solver.cpp:323] Optimization Done.
I0224 00:32:04.087548 655 caffe.cpp:222] Optimization Done.

```

If you look at output or the `logreg_auto_train.prototxt`, you'll see that the model is simple logistic regression. We can make it a little more advanced by introducing a non-linearity between weights that take the input and weights that give the output -- now we have a two-layer network. That

network is given in `nonlinear_auto_train.prototxt`, and that's the only change made in `nonlinear_logreg_solver.prototxt` which we will now use.

The final accuracy of the new network should be higher than logistic regression!

In [9]:

```
from caffe import layers as L
from caffe import params as P

def nonlinear_net(hdf5, batch_size):
    # one small nonlinearity, one leap for model kind
    n = caffe.NetSpec()
    n.data, n.label = L.HDF5Data(batch_size=batch_size, source=hdf5, ntop=2)
    # define a hidden layer of dimension 40
    n.ip1 = L.InnerProduct(n.data, num_output=40, weight_filler=dict(type='xavier'))
    # transform the output through the ReLU (rectified linear) non-linearity
    n.relu1 = L.ReLU(n.ip1, in_place=True)
    # score the (now non-linear) features
    n.ip2 = L.InnerProduct(n.relu1, num_output=2, weight_filler=dict(type='xavier'))
    # same accuracy and loss as before
    n.accuracy = L.Accuracy(n.ip2, n.label)
    n.loss = L.SoftmaxWithLoss(n.ip2, n.label)
    return n.to_proto()

train_net_path = 'examples/hdf5_classification/nonlinear_auto_train.prototxt'
with open(train_net_path, 'w') as f:
    f.write(str(nonlinear_net('examples/hdf5_classification/data/train.txt', 10)))

test_net_path = 'examples/hdf5_classification/nonlinear_auto_test.prototxt'
with open(test_net_path, 'w') as f:
    f.write(str(nonlinear_net('examples/hdf5_classification/data/test.txt', 10)))

solver_path = 'examples/hdf5_classification/nonlinear_logreg_solver.prototxt'
with open(solver_path, 'w') as f:
    f.write(str(solver(train_net_path, test_net_path)))
```

In [10]:

```
%%timeit
caffe.set_mode_cpu()
solver = caffe.get_solver(solver_path)
solver.solve()

accuracy = 0
batch_size = solver.test_nets[0].blobs['data'].num
test_iters = int(len(Xt) / batch_size)
for i in range(test_iters):
    solver.test_nets[0].forward()
    accuracy += solver.test_nets[0].blobs['accuracy'].data
accuracy /= test_iters

print("Accuracy: {:.3f}".format(accuracy))
```

Accuracy: 0.838

Accuracy: 0.837

Accuracy: 0.838

Accuracy: 0.834

1 loop, best of 3: 277 ms per loop

Do the same through the command line interface for detailed output on the model and solving.

In [11]:

```
!./build/tools/caffe train -solver examples/hdf5_classification/nonlinear_logreg_solver.prototxt
```

```
I0224 00:32:05.654265    658 caffe.cpp:178] Use CPU.
I0224 00:32:05.810444    658 solver.cpp:48] Initializing solver from parameters:
train_net: "examples/hdf5_classification/nonlinear_auto_train.prototxt"
test_net: "examples/hdf5_classification/nonlinear_auto_test.prototxt"
test_iter: 250
test_interval: 1000
base_lr: 0.01
display: 1000
max_iter: 10000
lr_policy: "step"
gamma: 0.1
momentum: 0.9
weight_decay: 0.0005
stepsize: 5000
snapshot: 10000
snapshot_prefix: "examples/hdf5_classification/data/train"
solver_mode: CPU
I0224 00:32:05.810634    658 solver.cpp:81] Creating training net from train_net file: examples/hdf5_classification/nonlinear_auto_train.prototxt
I0224 00:32:05.810835    658 net.cpp:49] Initializing net from parameters:
state {
  phase: TRAIN
}
layer {
  name: "data"
  type: "HDF5Data"
  top: "data"
  top: "label"
  hdf5_data_param {
    source: "examples/hdf5_classification/data/train.txt"
    batch_size: 10
  }
}
layer {
  name: "ip1"
  type: "InnerProduct"
  bottom: "data"
  top: "ip1"
  inner_product_param {
    num_output: 40
    weight_filler {
      type: "xavier"
    }
  }
}
layer {
  name: "relu1"
  type: "ReLU"
  bottom: "ip1"
  top: "ip1"
}
```



```

layer {
  name: "ip2"
  type: "InnerProduct"
  bottom: "ip1"
  top: "ip2"
  inner_product_param {
    num_output: 2
    weight_filler {
      type: "xavier"
    }
  }
}
layer {
  name: "accuracy"
  type: "Accuracy"
  bottom: "ip2"
  bottom: "label"
  top: "accuracy"
}
layer {
  name: "loss"
  type: "SoftmaxWithLoss"
  bottom: "ip2"
  bottom: "label"
  top: "loss"
}

```

```

I0224 00:32:05.811061 658 layer_factory.hpp:77] Creating layer data
I0224 00:32:05.811079 658 net.cpp:106] Creating Layer data
I0224 00:32:05.811092 658 net.cpp:411] data -> data
I0224 00:32:05.811121 658 net.cpp:411] data -> label
I0224 00:32:05.811143 658 hdf5_data_layer.cpp:79] Loading list of HDF5 fi
lenames from: examples/hdf5_classification/data/train.txt
I0224 00:32:05.811189 658 hdf5_data_layer.cpp:93] Number of HDF5 files: 2
I0224 00:32:05.812254 658 hdf5.cpp:32] Datatype class: H5T_FLOAT
I0224 00:32:05.812677 658 net.cpp:150] Setting up data
I0224 00:32:05.812705 658 net.cpp:157] Top shape: 10 4 (40)
I0224 00:32:05.812721 658 net.cpp:157] Top shape: 10 (10)
I0224 00:32:05.812729 658 net.cpp:165] Memory required for data: 200
I0224 00:32:05.812739 658 layer_factory.hpp:77] Creating layer label_data
_l_split
I0224 00:32:05.812752 658 net.cpp:106] Creating Layer label_data_l_split
I0224 00:32:05.812762 658 net.cpp:454] label_data_l_split <- label
I0224 00:32:05.812774 658 net.cpp:411] label_data_l_split -> label_data_l
_split_0
I0224 00:32:05.812785 658 net.cpp:411] label_data_l_split -> label_data_l
_split_1
I0224 00:32:05.812798 658 net.cpp:150] Setting up label_data_l_split
I0224 00:32:05.812808 658 net.cpp:157] Top shape: 10 (10)
I0224 00:32:05.812816 658 net.cpp:157] Top shape: 10 (10)
I0224 00:32:05.812824 658 net.cpp:165] Memory required for data: 280
I0224 00:32:05.812831 658 layer_factory.hpp:77] Creating layer ip1
I0224 00:32:05.812841 658 net.cpp:106] Creating Layer ip1
I0224 00:32:05.812849 658 net.cpp:454] ip1 <- data
I0224 00:32:05.812860 658 net.cpp:411] ip1 -> ip1
I0224 00:32:05.813179 658 net.cpp:150] Setting up ip1
I0224 00:32:05.813196 658 net.cpp:157] Top shape: 10 40 (400)

I0224 00:32:05.813210 658 net.cpp:165] Memory required for data: 1880

```

```

I0224 00:32:05.813230 658 layer_factory.hpp:77] Creating layer relu1
I0224 00:32:05.813241 658 net.cpp:106] Creating Layer relu1
I0224 00:32:05.813251 658 net.cpp:454] relu1 <- ip1
I0224 00:32:05.813258 658 net.cpp:397] relu1 -> ip1 (in-place)
I0224 00:32:05.813271 658 net.cpp:150] Setting up relu1
I0224 00:32:05.813279 658 net.cpp:157] Top shape: 10 40 (400)
I0224 00:32:05.813287 658 net.cpp:165] Memory required for data: 3480
I0224 00:32:05.813294 658 layer_factory.hpp:77] Creating layer ip2
I0224 00:32:05.813304 658 net.cpp:106] Creating Layer ip2
I0224 00:32:05.813313 658 net.cpp:454] ip2 <- ip1
I0224 00:32:05.813321 658 net.cpp:411] ip2 -> ip2
I0224 00:32:05.813336 658 net.cpp:150] Setting up ip2
I0224 00:32:05.813345 658 net.cpp:157] Top shape: 10 2 (20)
I0224 00:32:05.813379 658 net.cpp:165] Memory required for data: 3560
I0224 00:32:05.813401 658 layer_factory.hpp:77] Creating layer ip2_ip2_0_
split
I0224 00:32:05.813417 658 net.cpp:106] Creating Layer ip2_ip2_0_split
I0224 00:32:05.813426 658 net.cpp:454] ip2_ip2_0_split <- ip2
I0224 00:32:05.813434 658 net.cpp:411] ip2_ip2_0_split -> ip2_ip2_0_split
_0
I0224 00:32:05.813446 658 net.cpp:411] ip2_ip2_0_split -> ip2_ip2_0_split
_1
I0224 00:32:05.813457 658 net.cpp:150] Setting up ip2_ip2_0_split
I0224 00:32:05.813465 658 net.cpp:157] Top shape: 10 2 (20)
I0224 00:32:05.813473 658 net.cpp:157] Top shape: 10 2 (20)
I0224 00:32:05.813480 658 net.cpp:165] Memory required for data: 3720
I0224 00:32:05.813488 658 layer_factory.hpp:77] Creating layer accuracy
I0224 00:32:05.813499 658 net.cpp:106] Creating Layer accuracy
I0224 00:32:05.813508 658 net.cpp:454] accuracy <- ip2_ip2_0_split_0
I0224 00:32:05.813515 658 net.cpp:454] accuracy <- label_data_1_split_0
I0224 00:32:05.813524 658 net.cpp:411] accuracy -> accuracy
I0224 00:32:05.813539 658 net.cpp:150] Setting up accuracy
I0224 00:32:05.813547 658 net.cpp:157] Top shape: (1)
I0224 00:32:05.813555 658 net.cpp:165] Memory required for data: 3724
I0224 00:32:05.813565 658 layer_factory.hpp:77] Creating layer loss
I0224 00:32:05.813585 658 net.cpp:106] Creating Layer loss
I0224 00:32:05.813599 658 net.cpp:454] loss <- ip2_ip2_0_split_1
I0224 00:32:05.813616 658 net.cpp:454] loss <- label_data_1_split_1
I0224 00:32:05.813627 658 net.cpp:411] loss -> loss
I0224 00:32:05.813642 658 layer_factory.hpp:77] Creating layer loss
I0224 00:32:05.813663 658 net.cpp:150] Setting up loss
I0224 00:32:05.813671 658 net.cpp:157] Top shape: (1)
I0224 00:32:05.813679 658 net.cpp:160] with loss weight 1
I0224 00:32:05.813695 658 net.cpp:165] Memory required for data: 3728
I0224 00:32:05.813704 658 net.cpp:226] loss needs backward computation.
I0224 00:32:05.813712 658 net.cpp:228] accuracy does not need backward co
mputation.
I0224 00:32:05.813720 658 net.cpp:226] ip2_ip2_0_split needs backward com
putation.
I0224 00:32:05.813729 658 net.cpp:226] ip2 needs backward computation.
I0224 00:32:05.813735 658 net.cpp:226] relu1 needs backward computation.
I0224 00:32:05.813743 658 net.cpp:226] ip1 needs backward computation.
I0224 00:32:05.813751 658 net.cpp:228] label_data_1_split does not need b
ackward computation.
I0224 00:32:05.813760 658 net.cpp:228] data does not need backward comput
ation.

I0224 00:32:05.813772 658 net.cpp:270] This network produces output accur

```

```
acy
I0224 00:32:05.813787 658 net.cpp:270] This network produces output loss
I0224 00:32:05.813809 658 net.cpp:283] Network initialization done.
I0224 00:32:05.813905 658 solver.cpp:181] Creating test net (#0) specified by test_net file: examples/hdf5_classification/nonlinear_auto_test.prototxt
I0224 00:32:05.813944 658 net.cpp:49] Initializing net from parameters:
state {
  phase: TEST
}
layer {
  name: "data"
  type: "HDF5Data"
  top: "data"
  top: "label"
  hdf5_data_param {
    source: "examples/hdf5_classification/data/test.txt"
    batch_size: 10
  }
}
layer {
  name: "ip1"
  type: "InnerProduct"
  bottom: "data"
  top: "ip1"
  inner_product_param {
    num_output: 40
    weight_filler {
      type: "xavier"
    }
  }
}
layer {
  name: "relu1"
  type: "ReLU"
  bottom: "ip1"
  top: "ip1"
}
layer {
  name: "ip2"
  type: "InnerProduct"
  bottom: "ip1"
  top: "ip2"
  inner_product_param {
    num_output: 2
    weight_filler {
      type: "xavier"
    }
  }
}
layer {
  name: "accuracy"
  type: "Accuracy"
  bottom: "ip2"
  bottom: "label"
  top: "accuracy"
}
```

```

layer {
  name: "loss"
  type: "SoftmaxWithLoss"
  bottom: "ip2"
  bottom: "label"
  top: "loss"
}

```

```

I0224 00:32:05.814131 658 layer_factory.hpp:77] Creating layer data
I0224 00:32:05.814142 658 net.cpp:106] Creating Layer data
I0224 00:32:05.814152 658 net.cpp:411] data -> data
I0224 00:32:05.814162 658 net.cpp:411] data -> label
I0224 00:32:05.814180 658 hdf5_data_layer.cpp:79] Loading list of HDF5 fi
lenames from: examples/hdf5_classification/data/test.txt
I0224 00:32:05.814220 658 hdf5_data_layer.cpp:93] Number of HDF5 files: 1
I0224 00:32:05.815207 658 net.cpp:150] Setting up data
I0224 00:32:05.815227 658 net.cpp:157] Top shape: 10 4 (40)
I0224 00:32:05.815243 658 net.cpp:157] Top shape: 10 (10)
I0224 00:32:05.815253 658 net.cpp:165] Memory required for data: 200
I0224 00:32:05.815260 658 layer_factory.hpp:77] Creating layer label_data
_l_split
I0224 00:32:05.815270 658 net.cpp:106] Creating Layer label_data_l_split
I0224 00:32:05.815279 658 net.cpp:454] label_data_l_split <- label
I0224 00:32:05.815287 658 net.cpp:411] label_data_l_split -> label_data_l
_split_0
I0224 00:32:05.815299 658 net.cpp:411] label_data_l_split -> label_data_l
_split_1
I0224 00:32:05.815310 658 net.cpp:150] Setting up label_data_l_split
I0224 00:32:05.815318 658 net.cpp:157] Top shape: 10 (10)
I0224 00:32:05.815326 658 net.cpp:157] Top shape: 10 (10)
I0224 00:32:05.815335 658 net.cpp:165] Memory required for data: 280
I0224 00:32:05.815341 658 layer_factory.hpp:77] Creating layer ip1
I0224 00:32:05.815351 658 net.cpp:106] Creating Layer ip1
I0224 00:32:05.815358 658 net.cpp:454] ip1 <- data
I0224 00:32:05.815367 658 net.cpp:411] ip1 -> ip1
I0224 00:32:05.815383 658 net.cpp:150] Setting up ip1
I0224 00:32:05.815398 658 net.cpp:157] Top shape: 10 40 (400)
I0224 00:32:05.815413 658 net.cpp:165] Memory required for data: 1880
I0224 00:32:05.815435 658 layer_factory.hpp:77] Creating layer relul
I0224 00:32:05.815450 658 net.cpp:106] Creating Layer relul
I0224 00:32:05.815459 658 net.cpp:454] relul <- ip1
I0224 00:32:05.815469 658 net.cpp:397] relul -> ip1 (in-place)
I0224 00:32:05.815479 658 net.cpp:150] Setting up relul
I0224 00:32:05.815486 658 net.cpp:157] Top shape: 10 40 (400)
I0224 00:32:05.815495 658 net.cpp:165] Memory required for data: 3480
I0224 00:32:05.815501 658 layer_factory.hpp:77] Creating layer ip2
I0224 00:32:05.815510 658 net.cpp:106] Creating Layer ip2
I0224 00:32:05.815518 658 net.cpp:454] ip2 <- ip1
I0224 00:32:05.815527 658 net.cpp:411] ip2 -> ip2
I0224 00:32:05.815542 658 net.cpp:150] Setting up ip2
I0224 00:32:05.815551 658 net.cpp:157] Top shape: 10 2 (20)
I0224 00:32:05.815559 658 net.cpp:165] Memory required for data: 3560
I0224 00:32:05.815570 658 layer_factory.hpp:77] Creating layer ip2_ip2_0_
split
I0224 00:32:05.815579 658 net.cpp:106] Creating Layer ip2_ip2_0_split
I0224 00:32:05.815587 658 net.cpp:454] ip2_ip2_0_split <- ip2
I0224 00:32:05.815600 658 net.cpp:411] ip2_ip2_0_split -> ip2_ip2_0_split

```

\_0

```

I0224 00:32:05.815619 658 net.cpp:411] ip2_ip2_0_split -> ip2_ip2_0_split
_1
I0224 00:32:05.815640 658 net.cpp:150] Setting up ip2_ip2_0_split
I0224 00:32:05.815654 658 net.cpp:157] Top shape: 10 2 (20)
I0224 00:32:05.815662 658 net.cpp:157] Top shape: 10 2 (20)
I0224 00:32:05.815670 658 net.cpp:165] Memory required for data: 3720
I0224 00:32:05.815677 658 layer_factory.hpp:77] Creating layer accuracy
I0224 00:32:05.815685 658 net.cpp:106] Creating Layer accuracy
I0224 00:32:05.815693 658 net.cpp:454] accuracy <- ip2_ip2_0_split_0
I0224 00:32:05.815702 658 net.cpp:454] accuracy <- label_data_1_split_0
I0224 00:32:05.815711 658 net.cpp:411] accuracy -> accuracy
I0224 00:32:05.815722 658 net.cpp:150] Setting up accuracy
I0224 00:32:05.815732 658 net.cpp:157] Top shape: (1)
I0224 00:32:05.815738 658 net.cpp:165] Memory required for data: 3724
I0224 00:32:05.815747 658 layer_factory.hpp:77] Creating layer loss
I0224 00:32:05.815754 658 net.cpp:106] Creating Layer loss
I0224 00:32:05.815762 658 net.cpp:454] loss <- ip2_ip2_0_split_1
I0224 00:32:05.815770 658 net.cpp:454] loss <- label_data_1_split_1
I0224 00:32:05.815779 658 net.cpp:411] loss -> loss
I0224 00:32:05.815790 658 layer_factory.hpp:77] Creating layer loss
I0224 00:32:05.815811 658 net.cpp:150] Setting up loss
I0224 00:32:05.815829 658 net.cpp:157] Top shape: (1)
I0224 00:32:05.815843 658 net.cpp:160]         with loss weight 1
I0224 00:32:05.815867 658 net.cpp:165] Memory required for data: 3728
I0224 00:32:05.815876 658 net.cpp:226] loss needs backward computation.
I0224 00:32:05.815884 658 net.cpp:228] accuracy does not need backward co
mputation.
I0224 00:32:05.815892 658 net.cpp:226] ip2_ip2_0_split needs backward com
putation.
I0224 00:32:05.815901 658 net.cpp:226] ip2 needs backward computation.
I0224 00:32:05.815908 658 net.cpp:226] relu1 needs backward computation.
I0224 00:32:05.815915 658 net.cpp:226] ip1 needs backward computation.
I0224 00:32:05.815923 658 net.cpp:228] label_data_1_split does not need b
ackward computation.
I0224 00:32:05.815932 658 net.cpp:228] data does not need backward comput
ation.
I0224 00:32:05.815938 658 net.cpp:270] This network produces output accur
acy
I0224 00:32:05.815946 658 net.cpp:270] This network produces output loss
I0224 00:32:05.815958 658 net.cpp:283] Network initialization done.
I0224 00:32:05.815978 658 solver.cpp:60] Solver scaffolding done.
I0224 00:32:05.816000 658 caffe.cpp:219] Starting Optimization
I0224 00:32:05.816016 658 solver.cpp:280] Solving
I0224 00:32:05.816030 658 solver.cpp:281] Learning Rate Policy: step
I0224 00:32:05.816048 658 solver.cpp:338] Iteration 0, Testing net (#0)
I0224 00:32:05.831967 658 solver.cpp:406]         Test net output #0: accurac
y = 0.4464
I0224 00:32:05.832033 658 solver.cpp:406]         Test net output #1: loss =
0.909841 (* 1 = 0.909841 loss)
I0224 00:32:05.832186 658 solver.cpp:229] Iteration 0, loss = 0.798509
I0224 00:32:05.832218 658 solver.cpp:245]         Train net output #0: accura
cy = 0.6
I0224 00:32:05.832247 658 solver.cpp:245]         Train net output #1: loss =
0.798509 (* 1 = 0.798509 loss)
I0224 00:32:05.832281 658 sgd_solver.cpp:106] Iteration 0, lr = 0.01
I0224 00:32:05.859506 658 solver.cpp:338] Iteration 1000, Testing net (#0
)

```

```
I0224 00:32:05.862799 658 solver.cpp:406] Test net output #0: accurac
y = 0.8156
I0224 00:32:05.862818 658 solver.cpp:406] Test net output #1: loss =
0.44259 (* 1 = 0.44259 loss)
I0224 00:32:05.862853 658 solver.cpp:229] Iteration 1000, loss = 0.537015
I0224 00:32:05.862864 658 solver.cpp:245] Train net output #0: accura
cy = 0.7
I0224 00:32:05.862875 658 solver.cpp:245] Train net output #1: loss =
0.537015 (* 1 = 0.537015 loss)
I0224 00:32:05.862885 658 sgd_solver.cpp:106] Iteration 1000, lr = 0.01
I0224 00:32:05.883155 658 solver.cpp:338] Iteration 2000, Testing net (#0
)
I0224 00:32:05.886435 658 solver.cpp:406] Test net output #0: accurac
y = 0.8116
I0224 00:32:05.886451 658 solver.cpp:406] Test net output #1: loss =
0.434079 (* 1 = 0.434079 loss)
I0224 00:32:05.886484 658 solver.cpp:229] Iteration 2000, loss = 0.43109
I0224 00:32:05.886497 658 solver.cpp:245] Train net output #0: accura
cy = 0.9
I0224 00:32:05.886508 658 solver.cpp:245] Train net output #1: loss =
0.43109 (* 1 = 0.43109 loss)
I0224 00:32:05.886518 658 sgd_solver.cpp:106] Iteration 2000, lr = 0.01
I0224 00:32:05.907243 658 solver.cpp:338] Iteration 3000, Testing net (#0
)
I0224 00:32:05.910521 658 solver.cpp:406] Test net output #0: accurac
y = 0.8168
I0224 00:32:05.910537 658 solver.cpp:406] Test net output #1: loss =
0.425661 (* 1 = 0.425661 loss)
I0224 00:32:05.910905 658 solver.cpp:229] Iteration 3000, loss = 0.430245
I0224 00:32:05.910922 658 solver.cpp:245] Train net output #0: accura
cy = 0.7
I0224 00:32:05.910933 658 solver.cpp:245] Train net output #1: loss =
0.430245 (* 1 = 0.430245 loss)
I0224 00:32:05.910943 658 sgd_solver.cpp:106] Iteration 3000, lr = 0.01
I0224 00:32:05.931205 658 solver.cpp:338] Iteration 4000, Testing net (#0
)
I0224 00:32:05.934479 658 solver.cpp:406] Test net output #0: accurac
y = 0.8324
I0224 00:32:05.934496 658 solver.cpp:406] Test net output #1: loss =
0.404891 (* 1 = 0.404891 loss)
I0224 00:32:05.934530 658 solver.cpp:229] Iteration 4000, loss = 0.628955
I0224 00:32:05.934542 658 solver.cpp:245] Train net output #0: accura
cy = 0.7
I0224 00:32:05.934553 658 solver.cpp:245] Train net output #1: loss =
0.628955 (* 1 = 0.628955 loss)
I0224 00:32:05.934583 658 sgd_solver.cpp:106] Iteration 4000, lr = 0.01
I0224 00:32:05.955108 658 solver.cpp:338] Iteration 5000, Testing net (#0
)
I0224 00:32:05.958377 658 solver.cpp:406] Test net output #0: accurac
y = 0.8364
I0224 00:32:05.958395 658 solver.cpp:406] Test net output #1: loss =
0.404235 (* 1 = 0.404235 loss)
I0224 00:32:05.958432 658 solver.cpp:229] Iteration 5000, loss = 0.394939
I0224 00:32:05.958444 658 solver.cpp:245] Train net output #0: accura
cy = 0.9
I0224 00:32:05.958456 658 solver.cpp:245] Train net output #1: loss =
0.39494 (* 1 = 0.39494 loss)
```

```
I0224 00:32:05.958466 658 sgd_solver.cpp:106] Iteration 5000, lr = 0.001
I0224 00:32:05.978703 658 solver.cpp:338] Iteration 6000, Testing net (#0
)
I0224 00:32:05.981973 658 solver.cpp:406] Test net output #0: accurac
y = 0.838
I0224 00:32:05.981991 658 solver.cpp:406] Test net output #1: loss =
0.385743 (* 1 = 0.385743 loss)
I0224 00:32:05.982347 658 solver.cpp:229] Iteration 6000, loss = 0.411537
I0224 00:32:05.982362 658 solver.cpp:245] Train net output #0: accura
cy = 0.8
I0224 00:32:05.982373 658 solver.cpp:245] Train net output #1: loss =
0.411537 (* 1 = 0.411537 loss)
I0224 00:32:05.982383 658 sgd_solver.cpp:106] Iteration 6000, lr = 0.001
I0224 00:32:06.003015 658 solver.cpp:338] Iteration 7000, Testing net (#0
)
I0224 00:32:06.006283 658 solver.cpp:406] Test net output #0: accurac
y = 0.8388
I0224 00:32:06.006301 658 solver.cpp:406] Test net output #1: loss =
0.384648 (* 1 = 0.384648 loss)
I0224 00:32:06.006335 658 solver.cpp:229] Iteration 7000, loss = 0.521072
I0224 00:32:06.006347 658 solver.cpp:245] Train net output #0: accura
cy = 0.7
I0224 00:32:06.006358 658 solver.cpp:245] Train net output #1: loss =
0.521073 (* 1 = 0.521073 loss)
I0224 00:32:06.006368 658 sgd_solver.cpp:106] Iteration 7000, lr = 0.001
I0224 00:32:06.026715 658 solver.cpp:338] Iteration 8000, Testing net (#0
)
I0224 00:32:06.029965 658 solver.cpp:406] Test net output #0: accurac
y = 0.8404
I0224 00:32:06.029983 658 solver.cpp:406] Test net output #1: loss =
0.380889 (* 1 = 0.380889 loss)
I0224 00:32:06.030015 658 solver.cpp:229] Iteration 8000, loss = 0.329477
I0224 00:32:06.030028 658 solver.cpp:245] Train net output #0: accura
cy = 0.9
I0224 00:32:06.030040 658 solver.cpp:245] Train net output #1: loss =
0.329477 (* 1 = 0.329477 loss)
I0224 00:32:06.030048 658 sgd_solver.cpp:106] Iteration 8000, lr = 0.001
I0224 00:32:06.050626 658 solver.cpp:338] Iteration 9000, Testing net (#0
)
I0224 00:32:06.053889 658 solver.cpp:406] Test net output #0: accurac
y = 0.8376
I0224 00:32:06.053906 658 solver.cpp:406] Test net output #1: loss =
0.382756 (* 1 = 0.382756 loss)
I0224 00:32:06.054271 658 solver.cpp:229] Iteration 9000, loss = 0.412227
I0224 00:32:06.054291 658 solver.cpp:245] Train net output #0: accura
cy = 0.8
I0224 00:32:06.054314 658 solver.cpp:245] Train net output #1: loss =
0.412228 (* 1 = 0.412228 loss)
I0224 00:32:06.054337 658 sgd_solver.cpp:106] Iteration 9000, lr = 0.001
I0224 00:32:06.074646 658 solver.cpp:456] Snapshotting to binary proto fi
le examples/hdf5_classification/data/train_iter_10000.caffemodel
I0224 00:32:06.074808 658 sgd_solver.cpp:273] Snapshotting solver state t
o binary proto file examples/hdf5_classification/data/train_iter_10000.solv
erstate
I0224 00:32:06.074889 658 solver.cpp:318] Iteration 10000, loss = 0.53279
8
I0224 00:32:06.074906 658 solver.cpp:338] Iteration 10000, Testing net (#
```

```
0)
I0224 00:32:06.078208    658 solver.cpp:406]    Test net output #0: accurac
y = 0.8388
I0224 00:32:06.078225    658 solver.cpp:406]    Test net output #1: loss =
0.382042 (* 1 = 0.382042 loss)
I0224 00:32:06.078234    658 solver.cpp:323] Optimization Done.
I0224 00:32:06.078241    658 caffe.cpp:222] Optimization Done.
```

In [12]:

```
# Clean up (comment this out if you want to examine the hdf5_classification/data directo
ry).
shutil.rmtree(dirname)
```