

Introduction of Convolution Neural Network

Yan Li

CCL@ILC

Yan3.li@intel.com

2015/7/20

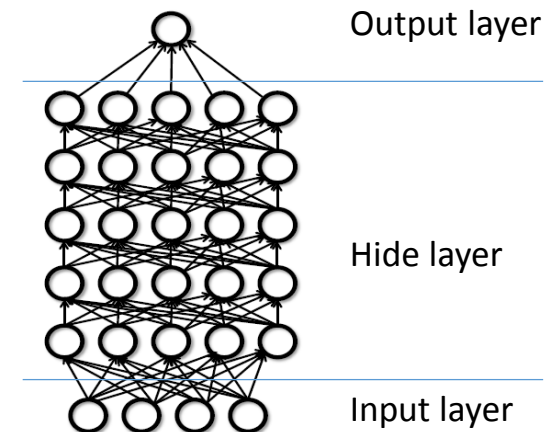
Deep Learning: Why so hot?

- Precondition: Big Data + HPC

"Deep architectures are compositions of many layers of adaptive non-linear components, in other words, they are cascades of parameterized non-linear modules that contain trainable parameters at all levels". ---- Bengio & Yann LeCun

- Deep networks are powerful than shallow ones
 - A highly flexible way to specify prior knowledge
 - More efficiently when the number of training examples becomes larger
 - Handle large families of functions, parameterized with millions of individual parameters
- Different deep architectures
 - Convolutional Neural Networks(CNNs)
 - Deep Belief Networks(DBNs)
- Open source

软件	核心开发语言	其它支持语言	CPU 支持	GPU 支持
Kaldi	C++/CUDA		✓	✓
Cuda-convnet	C++/CUDA	Python		✓
Caffe	C++/CUDA	Python, Matlab	✓	✓
Theano	Python		✓	✓
OverFeat	C++	Lua,Python	✓	
Torch7	Lua		✓	✓

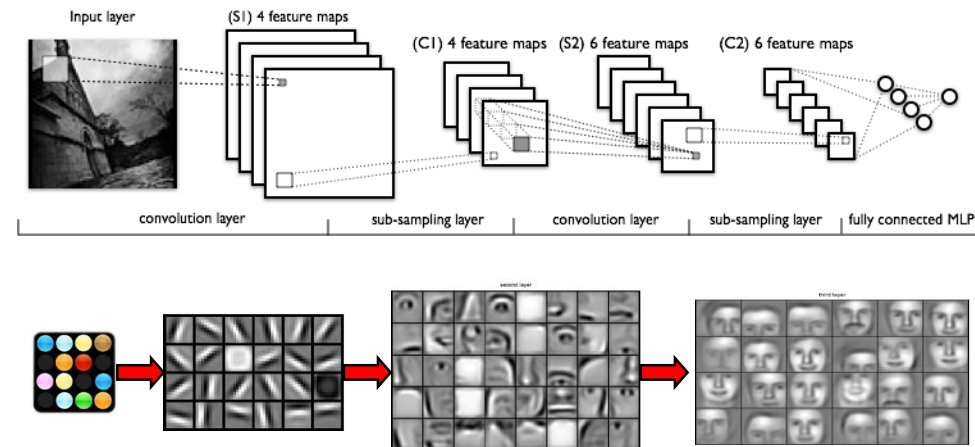
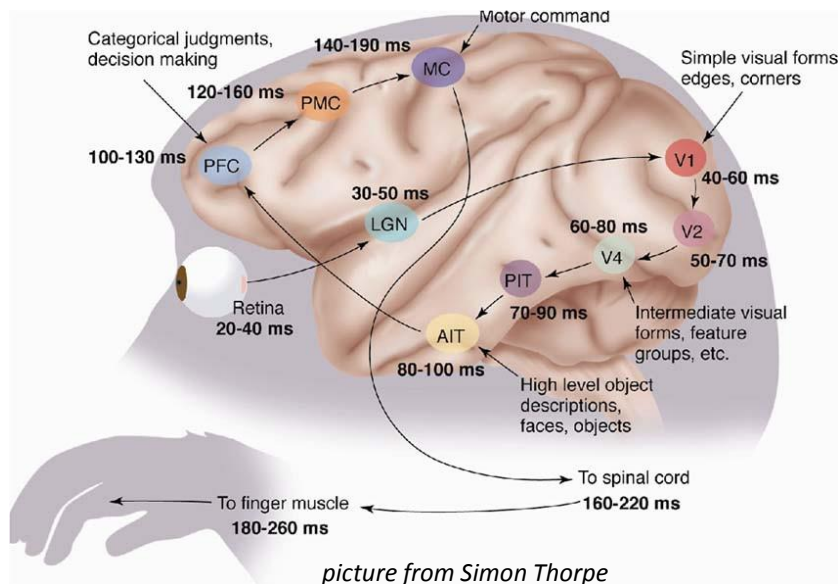


Deep Learning: Why so hot?

- Industry
 - Google: DistBelief
 - 1000 CPU nodes, each node has 16 cores
 - speech recognition & image recognition
 - Microsoft: Adam
 - 120 nodes, each node : dual Intel Xeon E5-2450L (16 cores)
 - Baidu: Minwa & Paddle
 - Speed & image recognition
 - Deep Image
 - Tencent: Mariana
 - Advertising in QQ and QQ Zone
 - Speed & image recognition in WeChat

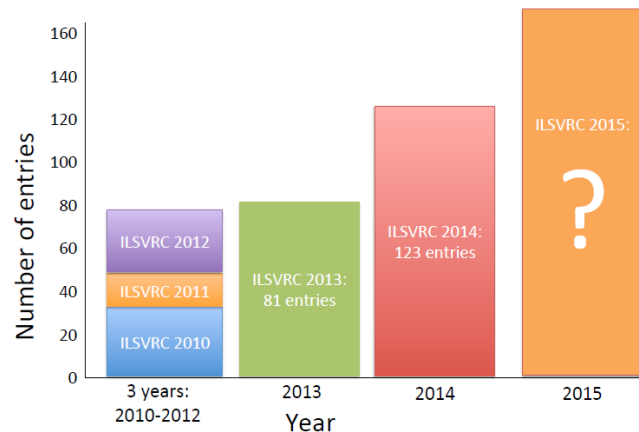
Deep Learning: Why so hot?

- Multiple layers work to build an improved feature space
- Highly varying functions can be efficiently represented with deep architectures
- Learning Representations / Features
 - Image recognition
 - Image -> low-level feature -> mid-level feature -> High-level feature -> Trainable Classifier
 - Pixel -> Edge -> texton -> motif -> part -> object
 - Text
 - Character->word->word group->clause->sentence->story
 - Speed
 - Sample->spectral band -> sound ->..->phone->phoneme->word



ImageNet Large Scale Visual Recognition Challenge

- A benchmark in object category classification and detection
 - <http://image-net.org/challenges/LSVRC/2015/>
- Participation in ILSVRC over years



ILSVRC overview: past, present, and future (<http://image-net.org/tutorials/cvpr2015/>)

Image classification annotations (1000 object classes)

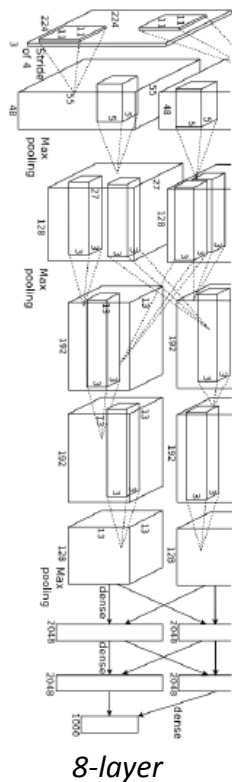
Year	Train images (per class)	Val images (per class)	Test images (per class)
ILSVRC2010	1,261,406 (668-3047)	50,000 (50)	150,000 (150)
ILSVRC2011	1,229,413 (384-1300)	50,000 (50)	100,000 (100)
ILSVRC2012-14	1,281,167 (732-1300)	50,000 (50)	100,000 (100)

- Image classification
 - Down to < 0.05 error (top-5 error) since ILSVRC2014
 - 2011: 0.26 -> 2012: 0.16
 - Deep learning method

ImageNet Large Scale Visual Recognition Challenge

Year 2012

SuperVision



[Krizhevsky NIPS 2012]

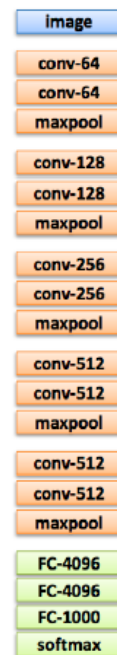
Year 2014

GoogLeNet



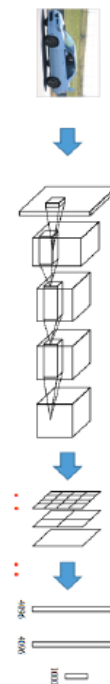
[Szegedy arxiv 2014]

VGG



16-layer

MSRA



22-layer

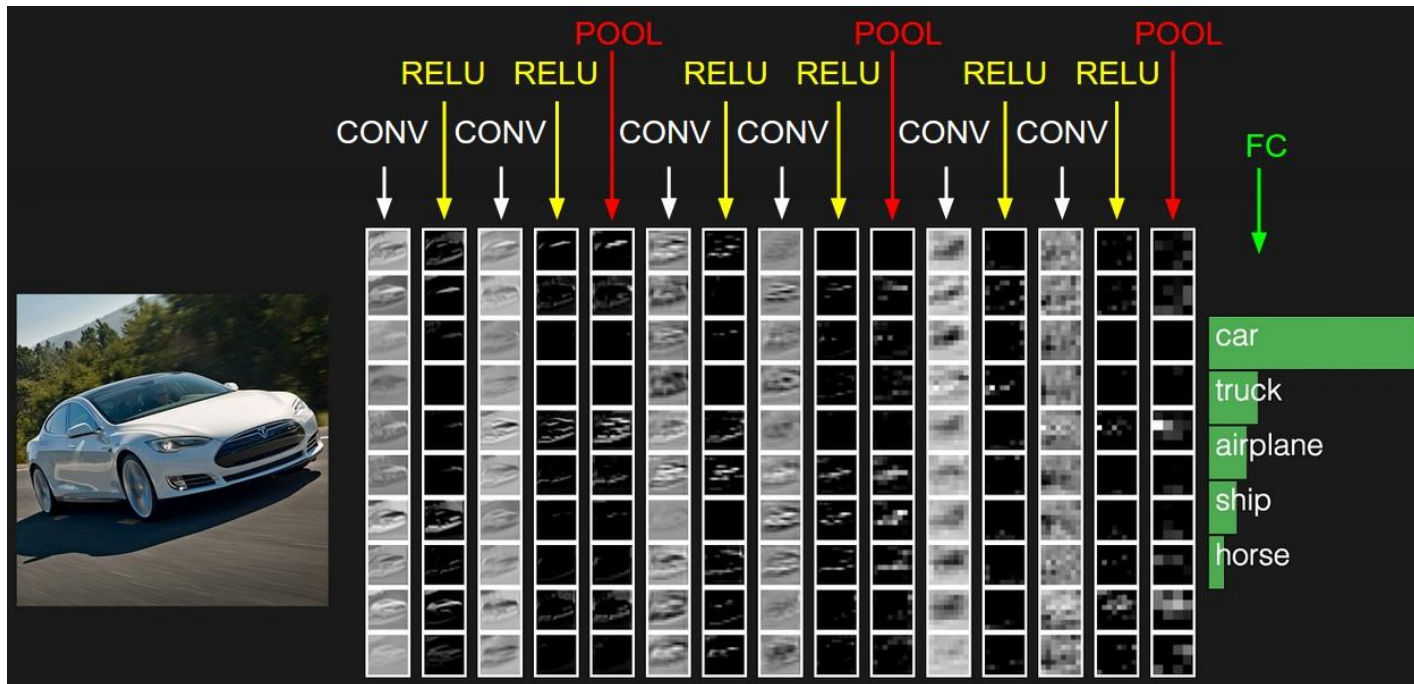
35/36 teams used
deep learning

20/36 teams used open-source Caffe implementation

ILSVRC overview: past, present, and future (<http://image-net.org/tutorials/cvpr2015/>)

The CNN Architecture

- The activations of an example ConvNet architecture

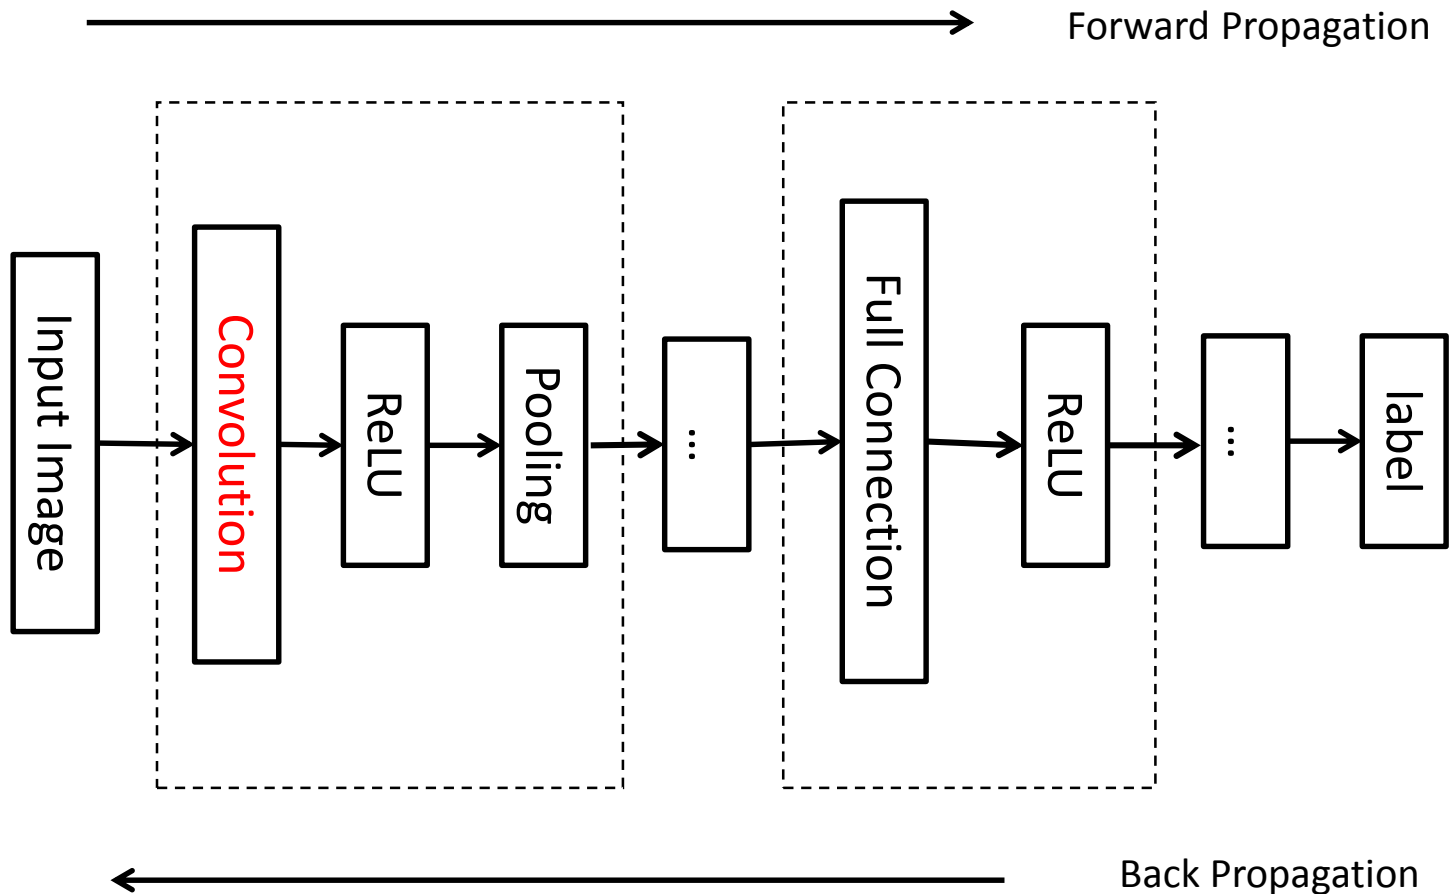


- Typical ConvNets

- [CONV-RELU-POOL] x N, [FC-RELU] x M, Softmax or
 - [CONV-RELU-CONV-RELU-POOL] x N, [FC-RELU] x M, FC, SOFTMAX
- $N \geq 0, M \geq 0$

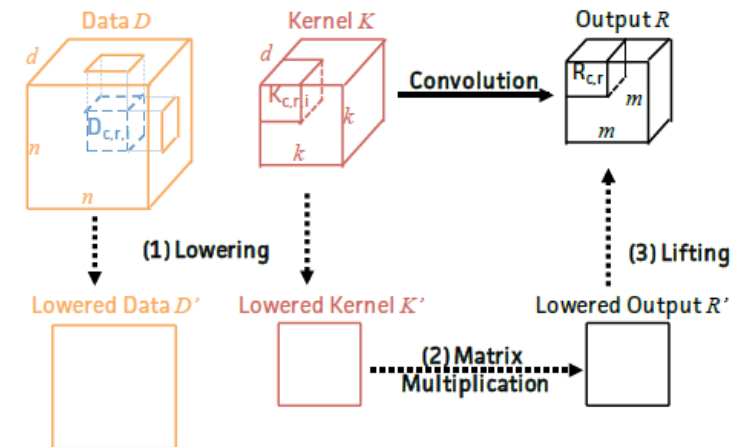
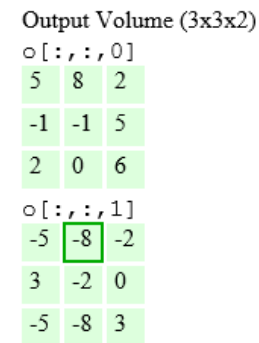
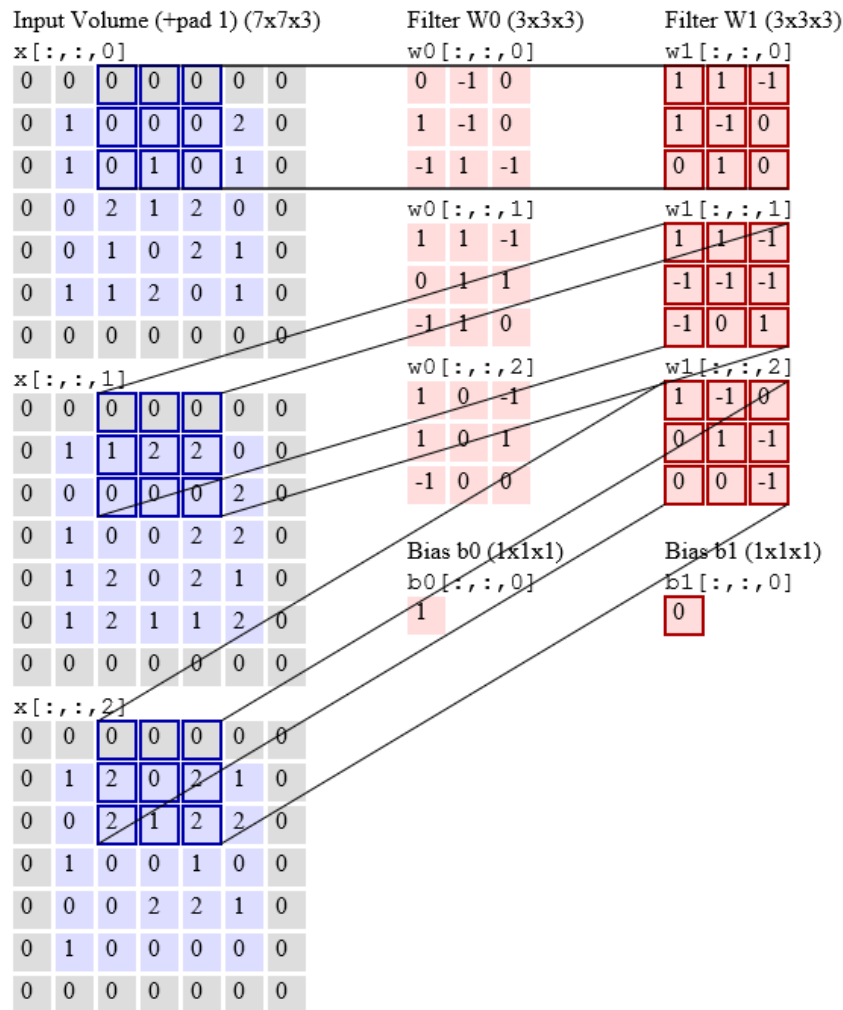
The CNN Architecture

- The classical computation flow for ConvNet



The CNN Architecture

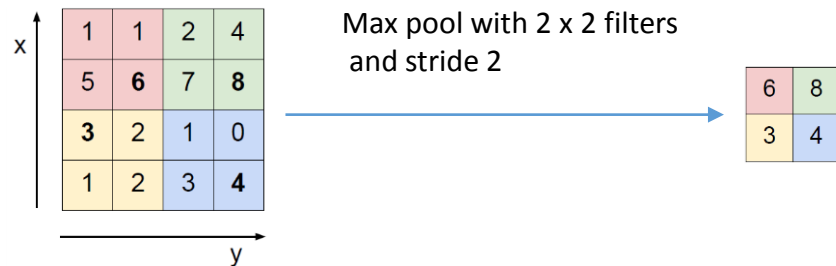
- The convolution layer



Convert to GEMM

The CNN Architecture

- CONV:
 - The Convolution layer
- RELU:
 - Non-linearity function
 - Rectifier function, such as $f(x) = \text{Max}(0, x)$
- POOL
 - downsamples



The CNN Architecture

- An example ConvNet architecture (VGG: 16 weight layers)

INPUT: [224x224x3] memory: $224*224*3=150K$ params: 0 (not counting biases)

CONV3-64: [224x224x64] memory: $224*224*64=3.2M$ params: $(3*3*3)*64 = 1,728$

CONV3-64: [224x224x64] memory: $224*224*64=3.2M$ params: $(3*3*64)*64 = 36,864$

POOL2: [112x112x64] memory: $112*112*64=800K$ params: 0

CONV3-128: [112x112x128] memory: $112*112*128=1.6M$ params: $(3*3*64)*128 = 73,728$

CONV3-128: [112x112x128] memory: $112*112*128=1.6M$ params: $(3*3*128)*128 = 147,456$

POOL2: [56x56x128] memory: $56*56*128=400K$ params: 0

CONV3-256: [56x56x256] memory: $56*56*256=800K$ params: $(3*3*128)*256 = 294,912$

CONV3-256: [56x56x256] memory: $56*56*256=800K$ params: $(3*3*256)*256 = 589,824$

CONV3-256: [56x56x256] memory: $56*56*256=800K$ params: $(3*3*256)*256 = 589,824$

POOL2: [28x28x256] memory: $28*28*256=200K$ params: 0

CONV3-512: [28x28x512] memory: $28*28*512=400K$ params: $(3*3*256)*512 = 1,179,648$

CONV3-512: [28x28x512] memory: $28*28*512=400K$ params: $(3*3*512)*512 = 2,359,296$

CONV3-512: [28x28x512] memory: $28*28*512=400K$ params: $(3*3*512)*512 = 2,359,296$

POOL2: [14x14x512] memory: $14*14*512=100K$ params: 0

CONV3-512: [14x14x512] memory: $14*14*512=100K$ params: $(3*3*512)*512 = 2,359,296$

CONV3-512: [14x14x512] memory: $14*14*512=100K$ params: $(3*3*512)*512 = 2,359,296$

CONV3-512: [14x14x512] memory: $14*14*512=100K$ params: $(3*3*512)*512 = 2,359,296$

POOL2: [7x7x512] memory: $7*7*512=25K$ params: 0

FC: [1x1x4096] memory: 4096 params: $7*7*512*4096 = 102,760,448$

FC: [1x1x4096] memory: 4096 params: $4096*4096 = 16,777,216$

FC: [1x1x1000] memory: 1000 params: $4096*1000 = 4,096,000$

Note:

Most memory is in
early CONV

Most params are
in late FC

TOTAL memory: $24M * 4 \text{ bytes} \approx 93MB$ / image (only forward! ~ 2 for bwd)

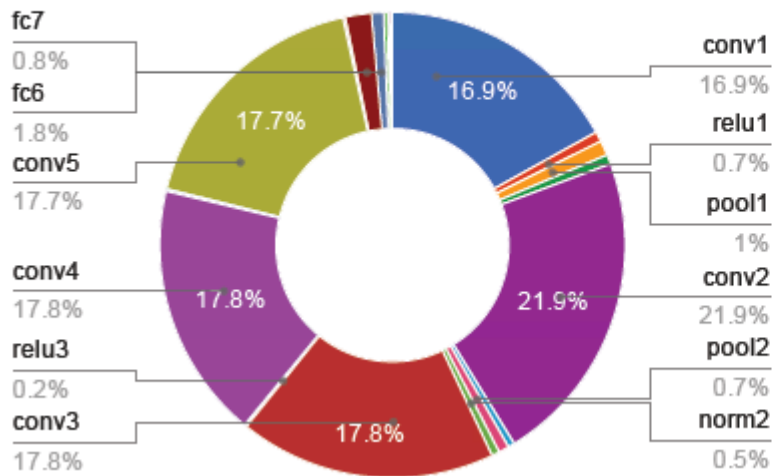
TOTAL params: 138M parameters

Fei-Fei Li & Andrej Karpathy. Visualizing and Understanding Convolution Neural Networks.

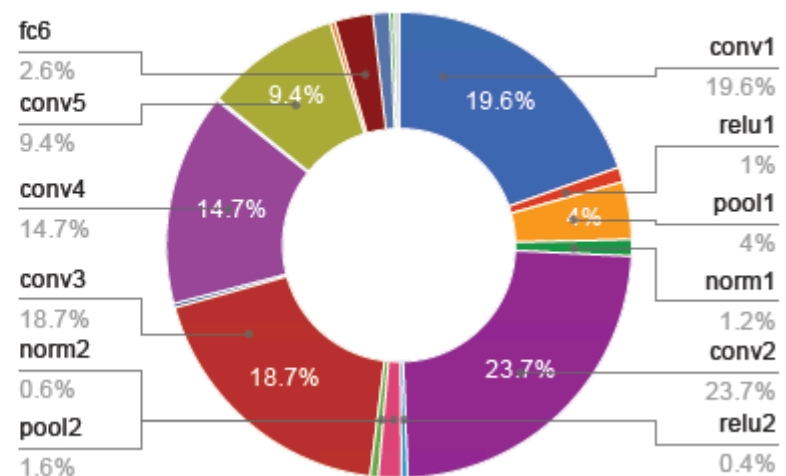
The CNN Architecture

- Computation time distribution of individual layers (Alex ConvNet)

GPU Forward Time Distribution



CPU Forward Time Distribution



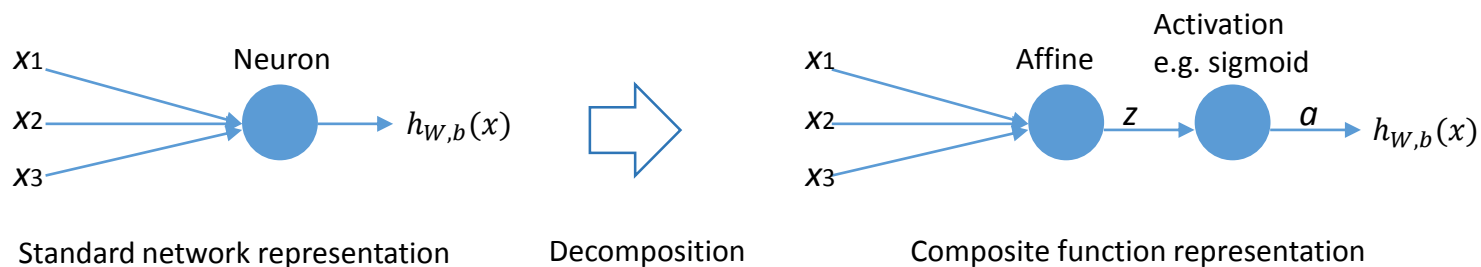
Yangqing Jia. "Learning Semantic Image Representations at a Large Scale"

Computation : $\text{Rate}_{\text{Conv}} > 80\%$

Communication : $\text{Rate}_{\text{FC}} > 80\%$

Backpropagation in CNN

- A neural network is decomposed into a composite function where each function element corresponds to a differentiable operation.
- Single neuron (the simplest neural network) example
 - A single neuron is decomposed into a composite function of an affine function element parameterized by \mathbf{W} and \mathbf{b} and an activation function element \mathbf{f} which we choose to be the sigmoid function.
 - Derivative of both affine and sigmoid function elements w.r.t. both inputs and parameters are known. Note that sigmoid function doesn't have neither parameters nor derivatives parameters.
 - Sigmoid function is applied element-wise. " \bullet " denotes Hadamard product, or element-wise product.



$$h_{W,b}(x) = f(W^T x + b) = \text{sigmoid}(\text{affine}_{W,b}(x)) = (\text{sigmoid} \circ \text{affine}_{W,b})(x)$$

$$\frac{\partial a}{\partial z} = a \bullet (1 - a) \text{ where } a = h_{W,b}(x) = \text{sigmoid}(z) = \frac{1}{1 + \exp(-z)}$$

$$\frac{\partial z}{\partial x} = W, \frac{\partial z}{\partial W} = x, \frac{\partial z}{\partial b} = I \text{ where } z = \text{affine}_{W,b}(x) = W^T x + b$$

Backpropagation in CNN

- Error signals are defined as the derivative of any cost function J which we choose to be the square error. Error signals are computed (propagated backward) by the chain rule of derivative and useful for computing the gradient of the cost function.

- Single neuron example

- Suppose we have m labeled training examples $\{(x^{(1)}, y^{(1)}), \dots, (x^{(m)}, y^{(m)})\}$. Square error cost function for each example is as follows. Overall cost function is the summation of cost functions over all examples.

$$J(W, b; x, y) = \frac{1}{2} \|y - h_{W,b}(x)\|^2$$

- Error signals of the square error cost function for each example are propagated using derivatives of function elements w.r.t. input.

$$\delta^{(a)} = \frac{\partial}{\partial a} J(W, b; x, y) = -(y - a)$$

$$\delta^{(z)} = \frac{\partial}{\partial z} J(W, b; x, y) = \frac{\partial J}{\partial a} \frac{\partial a}{\partial z} = \delta^{(a)} \cdot a \cdot (1 - a)$$

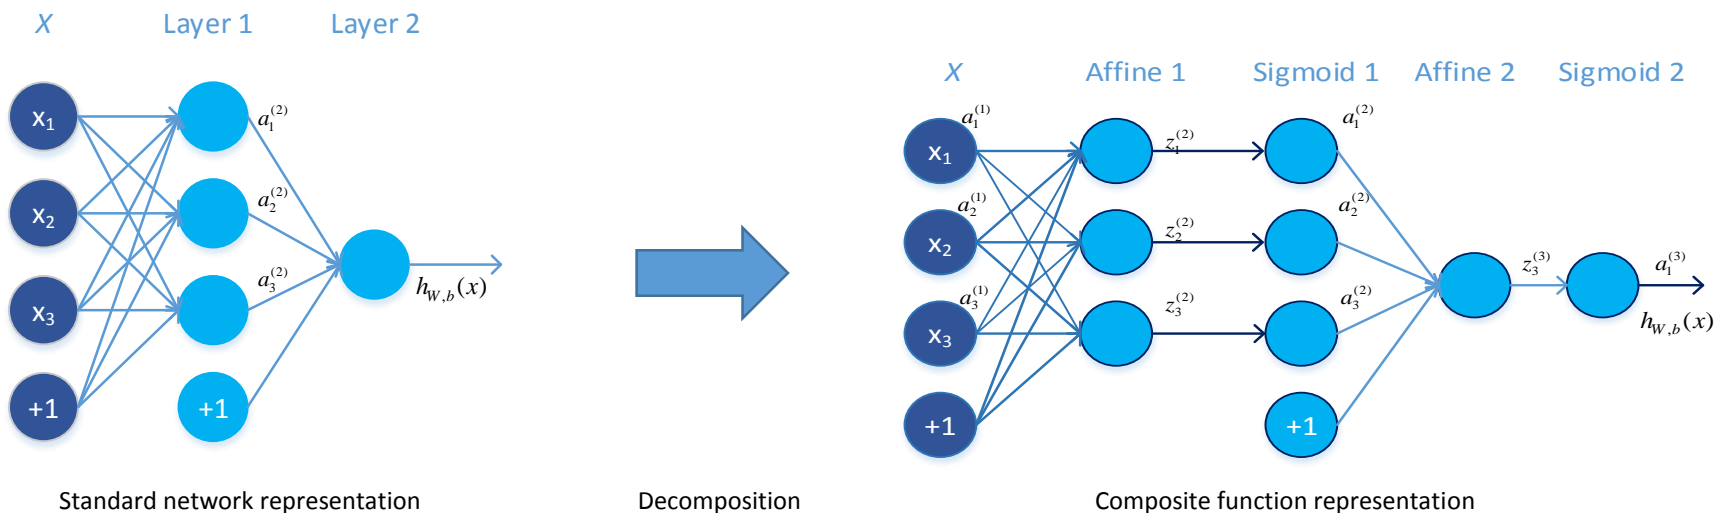
- Gradient of the cost function w.r.t parameters for each example is computed using error signals and derivatives of function elements w.r.t parameters. Summing gradients for all examples gets overall gradient.

$$\nabla_W J(W, b; x, y) = \frac{\partial}{\partial W} J(W, b; x, y) = \frac{\partial J}{\partial z} \frac{\partial z}{\partial W} = \delta^{(z)} x^T$$

$$\nabla_b J(W, b; x, y) = \frac{\partial}{\partial b} J(W, b; x, y) = \frac{\partial J}{\partial z} \frac{\partial z}{\partial b} = \delta^{(z)}$$

Backpropagation in CNN

- Composite function representation of a multi-layer neural network



$$h_{W,b}(x) = (\text{sigmoid} \circ \text{affine}_{W^{(2)},b^{(2)}} \circ \text{sigmoid} \circ \text{affine}_{W^{(1)},b^{(1)}})(x)$$

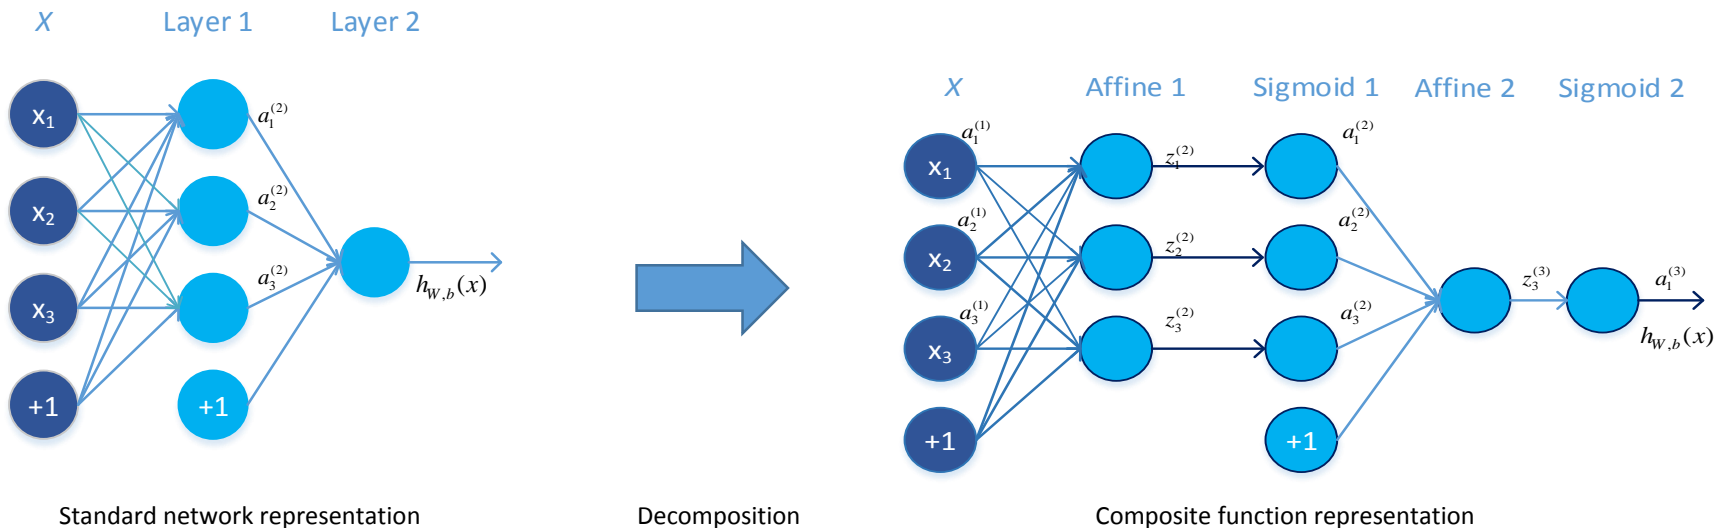
- Derivative of function elements w.r.t. inputs and parameters

$$a^{(1)} = x, a^{(l_{\max})} = h_{W,b}(x)$$

$$\frac{\partial a^{(l+1)}}{\partial z^{(l+1)}} = a^{(l+1)} \cdot (1 - a^{(l+1)}) \text{ where } a^{(l+1)} = \text{sigmoid}(z^{(l+1)}) = \frac{1}{1 + \exp(-z^{(l+1)})}$$

Backpropagation in CNN

- Error signals of the square error cost function for each example



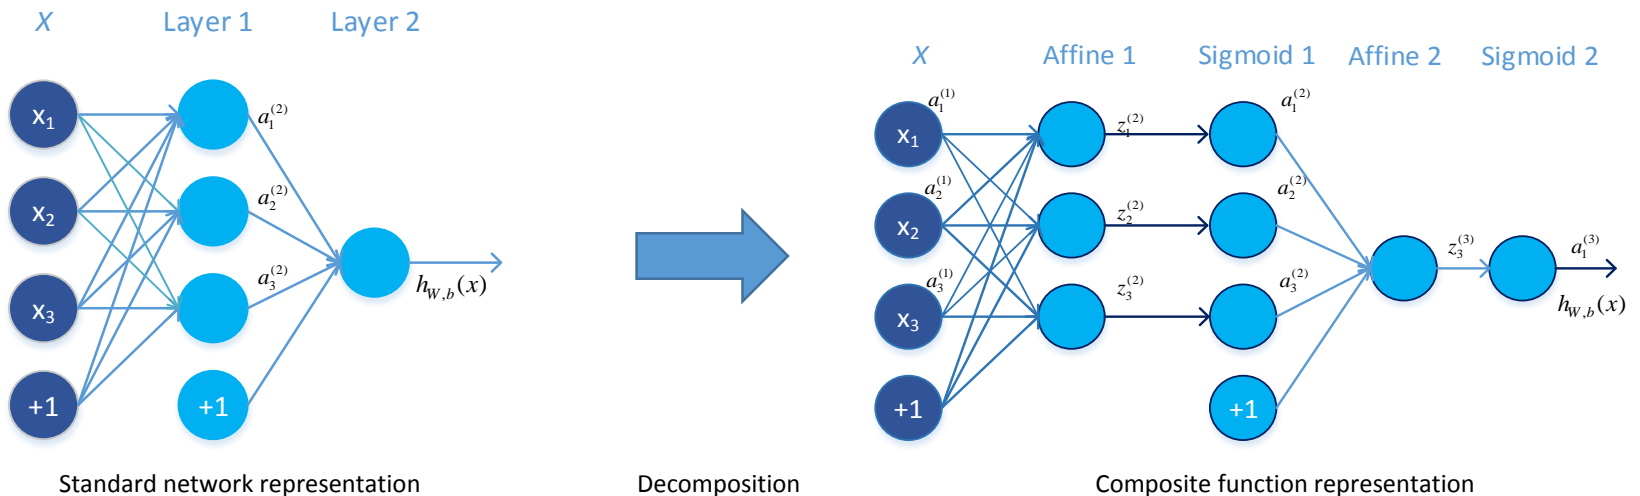
$$h_{W,b}(x) = (\text{sigmoid} \circ \text{affine}_{W^{(2)},b^{(2)}} \circ \text{sigmoid} \circ \text{affine}_{W^{(1)},b^{(1)}})(x)$$

$$\delta(a^{(l)}) = \frac{\partial}{\partial a^{(l)}} J(W, b; x, y) = \begin{cases} -(y - a^{(l)}) & \text{for } l = l_{\max} \\ \frac{\partial J}{\partial z^{(l+1)}} \frac{\partial z^{(l+1)}}{\partial a^{(l)}} = (W^{(l)})^T \delta(z^{(l+1)}) & \text{otherwise} \end{cases}$$

$$\delta(z^{(l)}) = \frac{\partial}{\partial z^{(l)}} J(W, b; x, y) = \frac{\partial J}{\partial a^{(l)}} \frac{\partial a^{(l)}}{\partial z^{(l)}} = \delta(a^{(l)}) \bullet a^{(l)} \bullet (1 - a^{(l)})$$

Backpropagation in CNN

- Gradient of the cost function w.r.t. parameters for each example



$$h_{W,b}(x) = (\text{sigmoid} \circ \text{affine}_{W^{(2)},b^{(2)}} \circ \text{sigmoid} \circ \text{affine}_{W^{(1)},b^{(1)}})(x)$$

$$\nabla_{W^{(l)}} J(W, b; x, y) = \frac{\partial}{\partial W^{(l)}} J(W, b; x, y) = \frac{\partial J}{\partial z^{(l+1)}} \frac{\partial z^{(l+1)}}{\partial W^{(l)}} = \delta^{(z^{(l+1)})} (a^{(l)})^T$$

$$\nabla_{b^{(l)}} J(W, b; x, y) = \frac{\partial}{\partial b^{(l)}} J(W, b; x, y) = \frac{\partial J}{\partial z^{(l+1)}} \frac{\partial z^{(l+1)}}{\partial b^{(l)}} = \delta^{(z^{(l+1)})}$$

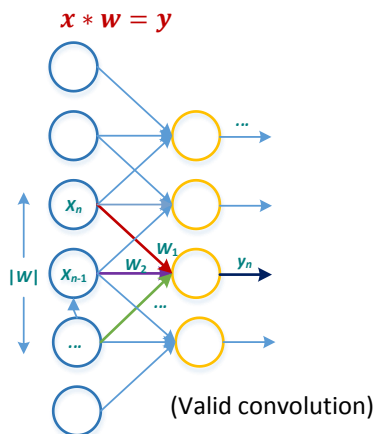
Derivative in Convolutional Layer

- Error signals and gradient for each example are computed by convolution using the commutativity property of convolution and the multivariable chain rule of derivative.
- Let's focus on single elements of error signals and a gradient w.r.t. w .

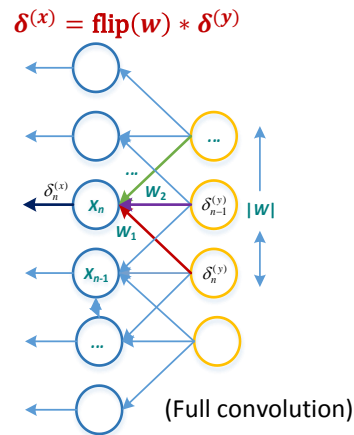
$$\delta_n^{(x)} = \frac{\partial J}{\partial x_n} = \frac{\partial J}{\partial y} \frac{\partial y}{\partial x_n} = \sum_{i=1}^{|w|} \frac{\partial J}{\partial y_{n-i+1}} \frac{\partial y_{n-i+1}}{\partial x_n} = \sum_{i=1}^{|w|} \delta_{n-i+1}^{(y)} w_i = (\delta^{(y)} * \text{flip}(w)) [n], \delta^{(x)} = [\delta_n^{(x)}] = \delta^{(y)} * \text{flip}(w)$$

Reverse order linear combination

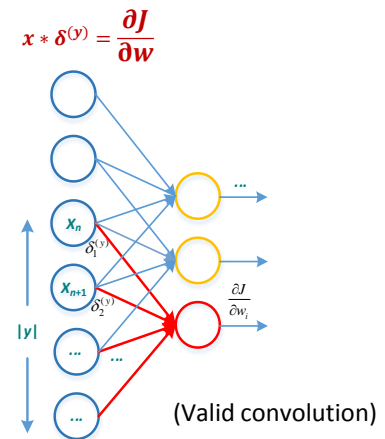
$$\frac{\partial J}{\partial w_i} = \frac{\partial J}{\partial y} \frac{\partial y}{\partial w_i} = \sum_{n=1}^{|x|-|w|+1} \frac{\partial J}{\partial y_n} \frac{\partial y_n}{\partial w_i} = \sum_{n=1}^{|x|-|w|+1} \delta_n^{(y)} x_{n+i-1} = (\delta^{(y)} * x) [i], \frac{\partial J}{\partial w} = \left[\frac{\partial y}{\partial w_i} \right] = \delta^{(y)} * x = x * \delta^{(y)}$$



Forward propagation (convolution)



backward propagation

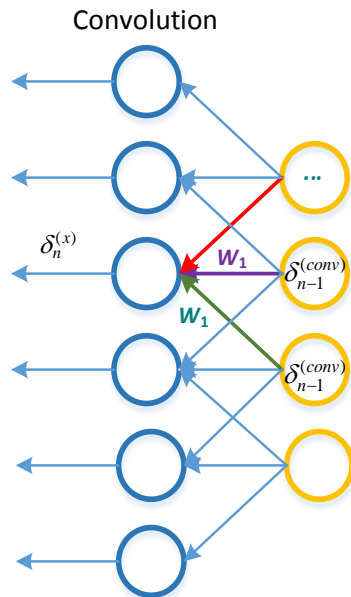


Gradient computation

Convolutional Neural Network

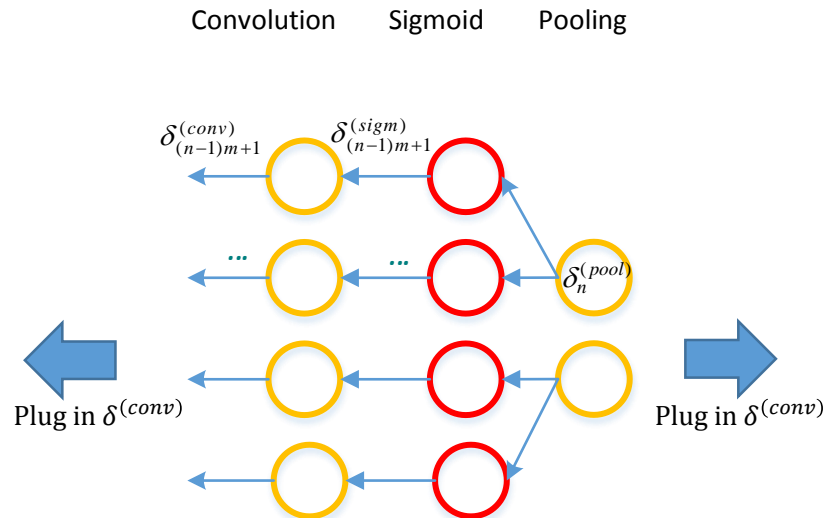
- Summary of Backpropagation

2. Propagate error signals $\delta^{(conv)}$



$$\delta^{(x)} = \delta^{(conv)} * \text{flip}(w)$$

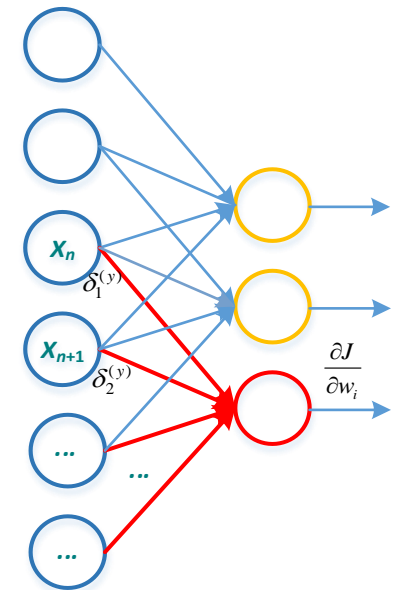
1. Propagate error signals $\delta^{(pool)}$



$$\delta^{(conv)} = \text{upsample}(\delta^{(conv)}, g') \cdot \boxed{f(\text{sigm}) \cdot (1 - f(\text{sigm}))}$$

Derivative of sigmoid

3. Compute gradient $\nabla_w J$



$$x * \delta^{(conv)} = \nabla_w J$$

Code Review: Full-connected Layer

- LayerSetUp()
 - Parameter initialization
 - For **one-time** initialization: reading parameters, fixed-size allocations, etc.
- Reshape()
 - For computing the sizes of top blobs, allocating buffers, and any other work that depends on the shapes of bottom blobs
- Differentiate
 - Reshape is called before every forward pass;
 - LayerSetUp is only called once at initialization. This allows networks to change their blob shapes while running.
 - <https://github.com/BVLC/caffe/issues/1385>

Code Review: Full-connected Layer

- GEMM: $C = \alpha A \times B + \beta C$

- A: $M \times K$
- B: $K \times N$
- C: $M \times N$

```
void caffe_cpu_gemm<float>(const CBLAS_TRANSPOSE TransA, const CBLAS_TRANSPOSE TransB,  
    const int M, const int N, const int K, const float alpha,  
    const float* A, const float* B, const float beta, float* C)
```

- GEMV: $Y = \alpha AX + \beta Y$

- A: $M \times N$
- X: $N \times 1$
- Y: $M \times 1$

```
void caffe_cpu_gemv<float>(const CBLAS_TRANSPOSE TransA, const int M,  
    const int N, const float alpha, const float* A, const float* x,  
    const float beta, float* y)
```

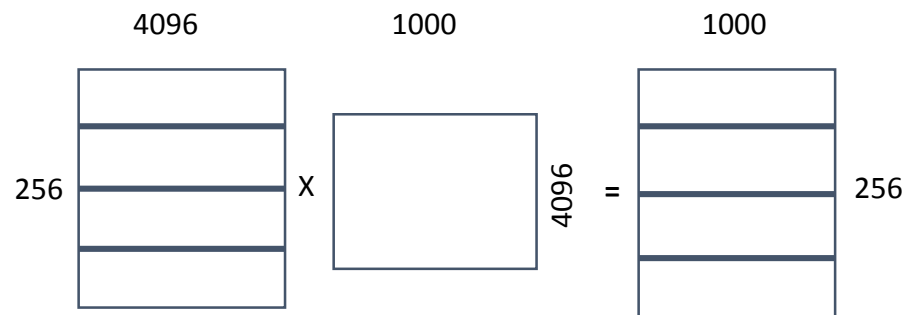
Code Review: Full-connected Layer

• Forward Computation

```
80 template <typename Dtype>
81 void InnerProductLayer<Dtype>::Forward_cpu(const vector<Blob<Dtype>*>& bottom,
82 const vector<Blob<Dtype>*>& top) {
83     const Dtype* bottom_data = bottom[0]->cpu_data();
84     Dtype* top_data = top[0]->mutable_cpu_data();
85     const Dtype* weight = this->blobs_[0]->cpu_data();
86     caffe_cpu_gemm<Dtype>(CblasNoTrans, CblasTrans, M_, N_, K_, (Dtype)1.,
87 bottom_data, weight, (Dtype)0., top_data); y=wx or y=xw'
88     if (bias_term_) {
89         caffe_cpu_gemm<Dtype>(CblasNoTrans, CblasNoTrans, M_, N_, 1, (Dtype)1.,
90 bias_multiplier_.cpu_data(), y=y+b
91 this->blobs_[1]->cpu_data(), (Dtype)1., top_data);
92     }
93 }
```

• Function: $y = wx + b$

- $M_$: the number of images
- $K_$: the number of features per image
- $N_$: the number of output neurons
 - X : $M \times K$ Y : $N \times 1$
 - W : $N \times K$ b : $N \times 1$

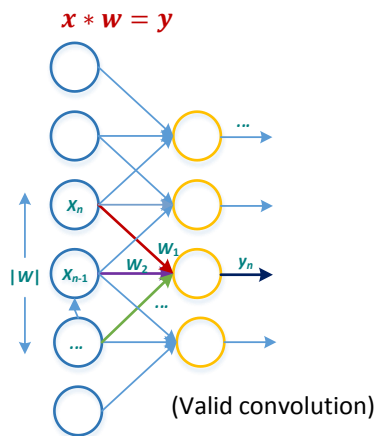


Code Review: Full-connected Layer

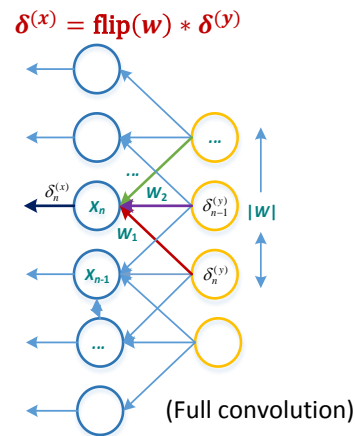
- Backward Computation

$$\nabla_W J(W, b; x, y) = \frac{\partial}{\partial W} J(W, b; x, y) = \frac{\partial J}{\partial z} \frac{\partial z}{\partial W} = \delta^{(z)} x^T$$

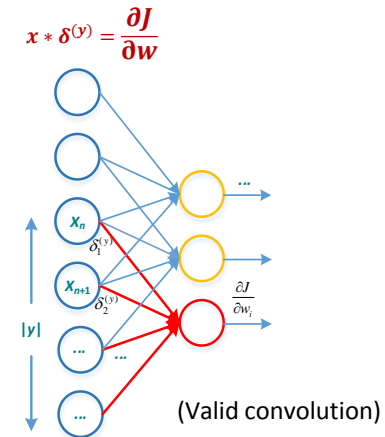
$$\nabla_b J(W, b; x, y) = \frac{\partial}{\partial b} J(W, b; x, y) = \frac{\partial J}{\partial y} \frac{\partial y}{\partial b} = \delta^{(y)}$$



Forward propagation (convolution)



backward propagation



Gradient computation

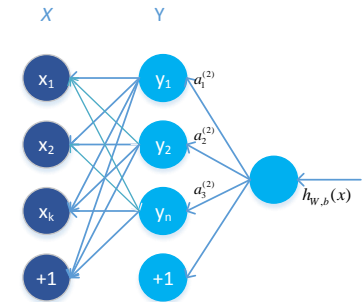
$$\delta_i^{(x)} = \frac{\partial J}{\partial x_i} = \frac{\partial J}{\partial y} \frac{\partial y}{\partial x_i} = \sum_{j=1}^{|w|} \frac{\partial J}{\partial y_{i-j+1}} \frac{\partial y_{i-j+1}}{\partial x_n} = \sum_{j=1}^{|w|} \delta_{i-j+1}^{(y)} w_j = (\delta^{(y)} * \text{flip}(w)) [n], \delta^{(x)} = [\delta_n^{(x)}] = \delta^{(y)} * \text{flip}(w)$$

Reverse order linear combination

$$\frac{\partial J}{\partial w_{ij}} = \frac{\partial J}{\partial y} \frac{\partial y}{\partial w_{ij}} = \frac{\partial J}{\partial y} \frac{\partial y}{\partial w_{ij}} = \delta_j^{(y)} * x_i \quad \frac{\partial J}{\partial b_i} = \delta_j^{(y)}$$

Code Review: Full-connected Layer

• Backward Computation



```

95 template <typename Dtype>
96 void InnerProductLayer<Dtype>::Backward_cpu(const vector<Blob<Dtype>*>& top,
97 const vector<bool>& propagate_down,
98 const vector<Blob<Dtype>*>& bottom) {
99     if (this->param_propagate_down_[0]) {
100         const Dtype* top_diff = top[0]->cpu_diff();
101         const Dtype* bottom_data = bottom[0]->cpu_data();
102         // Gradient with respect to weight
103         caffe_cpu_gemm<Dtype>(CblasTrans, CblasNoTrans, N_, K_, M_, (Dtype)1.,
104             top_diff, bottom_data, (Dtype)0., this->blobs_[0]->mutable_cpu_diff());
105     }
106     if (bias_term_ && this->param_propagate_down_[1]) {
107         const Dtype* top_diff = top[0]->cpu_diff();
108         // Gradient with respect to bias
109         caffe_cpu_gemv<Dtype>(CblasTrans, M_, N_, (Dtype)1., top_diff,
110             bias_multiplier_.cpu_data(), (Dtype)0.,
111             this->blobs_[1]->mutable_cpu_diff());
112     }
113     if (propagate_down[0]) {
114         const Dtype* top_diff = top[0]->cpu_diff();
115         // Gradient with respect to bottom data
116         caffe_cpu_gemm<Dtype>(CblasNoTrans, CblasNoTrans, M_, K_, N_, (Dtype)1.,
117             top_diff, this->blobs_[0]->cpu_data(), (Dtype)0.,
118             bottom[0]->mutable_cpu_diff());
119     }
120 }

```

$$\nabla_W J(W, b; x, y) = \frac{\partial}{\partial W} J(W, b; x, y) = \frac{\partial J}{\partial y} \frac{\partial y}{\partial W} = \delta^{(y)} x^T$$



$$\nabla_b J(W, b; x, y) = \frac{\partial}{\partial b} J(W, b; x, y) = \frac{\partial J}{\partial y} \frac{\partial y}{\partial b} = \delta^{(y)}$$

$$\delta_i^{(x)} = \left(\sum_{j=1}^{|y|} \delta_{i-j+1}^{(y)} w_{ji} \right) f'_{\theta_i}(z_i^{(x)})$$

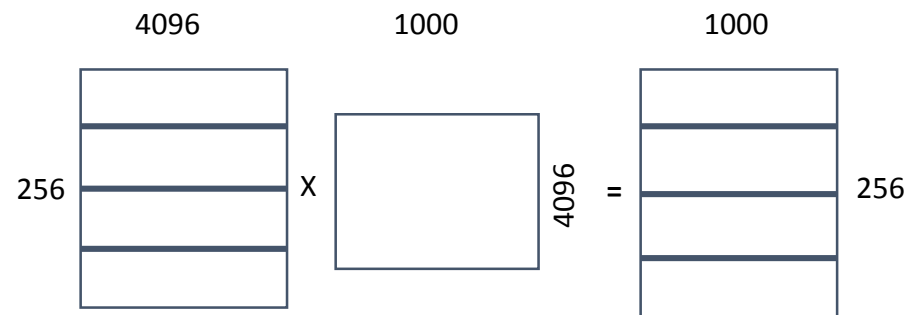
Code Review: Convolution Layer

• Forward Computation

```
80 template <typename Dtype>
81 void InnerProductLayer<Dtype>::Forward_cpu(const vector<Blob<Dtype>*>& bottom,
82 const vector<Blob<Dtype>*>& top) {
83     const Dtype* bottom_data = bottom[0]->cpu_data();
84     Dtype* top_data = top[0]->mutable_cpu_data();
85     const Dtype* weight = this->blobs_[0]->cpu_data();
86     caffe_cpu_gemm<Dtype>(CblasNoTrans, CblasTrans, M_, N_, K_, (Dtype)1.,
87 bottom_data, weight, (Dtype)0., top_data); y=wx or y=xw'
88     if (bias_term_) {
89         caffe_cpu_gemm<Dtype>(CblasNoTrans, CblasNoTrans, M_, N_, 1, (Dtype)1.,
90 bias_multiplier_.cpu_data(), y=y+b
91 this->blobs_[1]->cpu_data(), (Dtype)1., top_data);
92     }
93 }
```

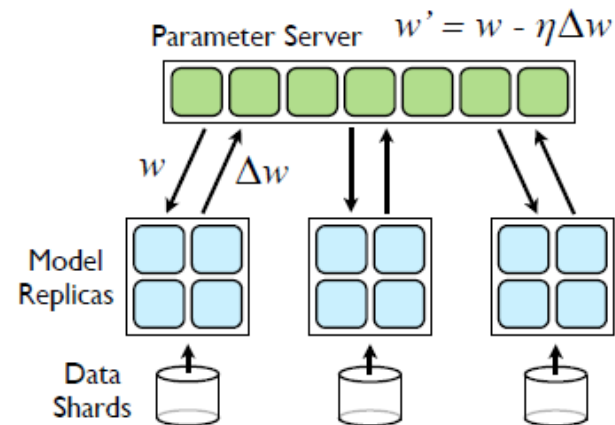
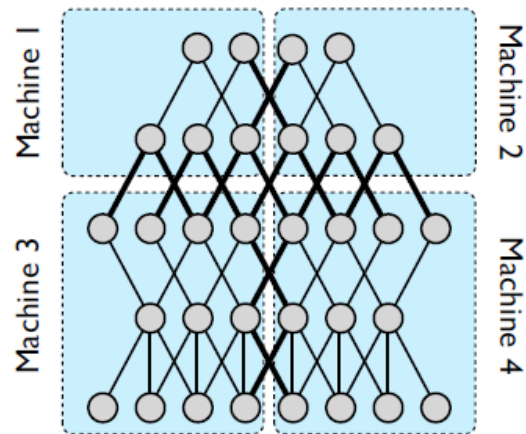
• Function: $y = wx + b$

- $M_$: the number of images
- $K_$: the number of features per image
- $N_$: the number of output neurons
 - X : $M \times K$ Y : $N \times 1$
 - W : $N \times K$ b : $N \times 1$



Distributed Deep Learning Networks

- Model Parallelism

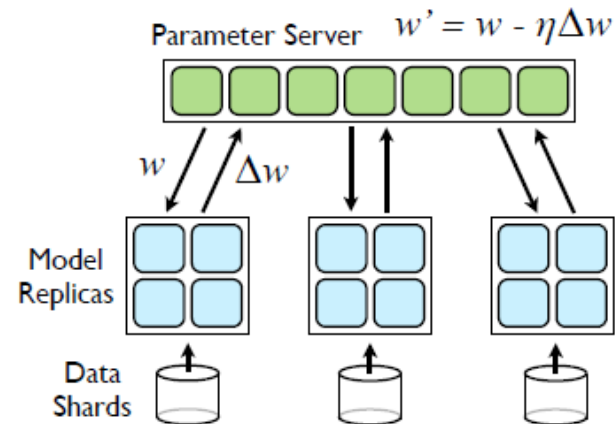
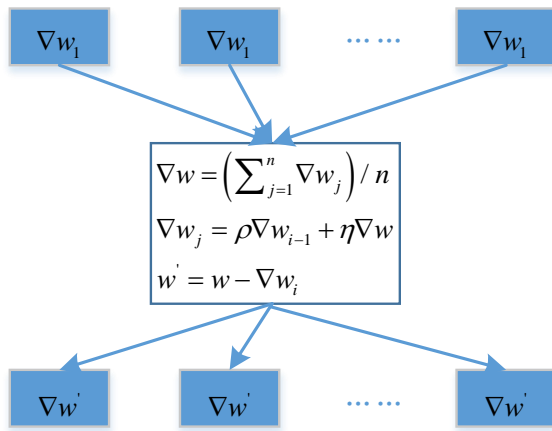


- Data Parallelism

- Based on mini-batch
- A single parameter server

Our Previous Work

- Data Parallelism



- Data Parallelism

- Based on mini-batch
- A single parameter server

- MPI + OpenMP

- MKL BIAS (GEMM & GEMV)

Distributed implementation

```
1 for i = 0; i < out_iter; i++ do
2   loss = net → ForwardBackward();           // 前向反向, 获得  $\nabla \mathbf{w}$  和 loss
3   ComputeUpdateValue();                     // 计算增量,  $\Delta \mathbf{w}_i = \rho \Delta \mathbf{w}_{i-1} + \eta \nabla \mathbf{w}$ 
4   net → Update();                           // 更新参数,  $\Delta \mathbf{w}' = \mathbf{w} - \Delta \mathbf{w}_i$ 
```

算法 1.1: Caffe 的串行训练过程伪代码

```
1 for i = 0; i < out_iter; i++ do
2   for j = 1 to n do in parallel
3     loss_j = net → ForwardBackward();       // 前向反向, 获得  $\nabla \mathbf{w}_j$  和 loss_j
4     GlobalSyncDiff();                       // 全局归约  $\nabla \mathbf{w}$  并求平均
5     GlobalSyncLoss();                       // 全局归约 loss 并求平均
6     if j == root then
7       ComputeUpdateValue();                 // 计算增量,  $\Delta \mathbf{w}_i = \rho \Delta \mathbf{w}_{i-1} + \eta \nabla \mathbf{w}$ 
8       net → Update();                       // 更新参数,  $\Delta \mathbf{w}' = \mathbf{w} - \Delta \mathbf{w}_i$ 
9       GlobalSyncData();                     // 广播新模型的参数  $\mathbf{w}'$ 
10
11 void GlobalSyncDiff()                      // length 是  $\mathbf{w}$  的维度
12   MPI_Reduce( $\nabla \mathbf{w}_j$ ,  $\nabla \mathbf{w}$ , length, MPI_FLOAT, MPI_SUM, root, COMM);
13   if j == root then
14     cblas_sscal(1/n,  $\nabla \mathbf{w}$ );              // 对  $\nabla \mathbf{w}$  求平均
15
16 void GlobalSyncLoss()
17   MPI_Reduce(loss_j, loss, 1, MPI_FLOAT, MPI_SUM, root, COMM);
18   if j == root then
19     loss = loss/n;
20
21 void GlobalSyncData()
22   MPI_Bcast( $\mathbf{w}'$ , length, MPI_FLOAT, root, COMM);
```

算法 1.2: Caffe 的分布式训练过程伪代码

Partial Synchronization

- 采用部分同步的非阻塞方式实现数据
 - **迭代内**：保持不变
 - **迭代间**：主节点只接收前 $s(s < n)$ 个返回的梯度，更新为最新的模型 w' 和 b' 后(更新公式不变)，也只广播给这 s 个节点
 - **缺点**：如果存在慢节点，它每次都不能在前 $s(s < n)$ 个返回梯度，那么这个节点将始终无法获得最新的模型参数
- 采用有界延迟策略保证训练过程的一致性
 - 每 t 次迭代就强制所有的节点同步更新参数

```
1 for i = 0; i < out_iter; i++ do
2   for j = 1 to n do in parallel
3     loss_j = net → ForwardBackward();           // 同算法1.2的第3行
4     GlobalAsyncDiff();      // 对前 s 个返回梯度的节点, 进行部分全局归约
5     GlobalSyncLoss();       // 计算 loss 值, 同算法1.2的第5行
6     if j == root then
7       ComputeUpdateValue();      // 计算增量, 同算法1.2的第7行
8       net → Update();            // 更新参数, 同算法1.2的第8行
9     GlobalAsyncData();          // 将新模型的参数  $w'$  广播给这 s 个节点
```

Evaluation

- 实验环境

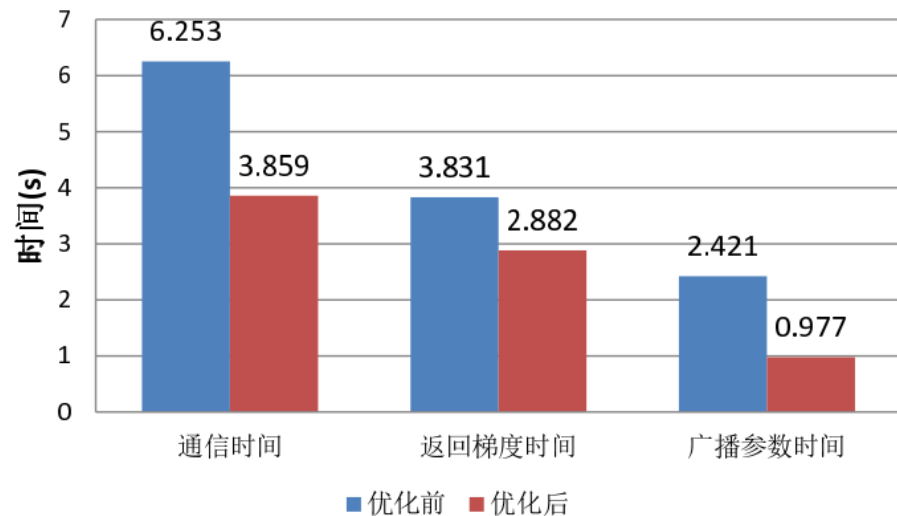
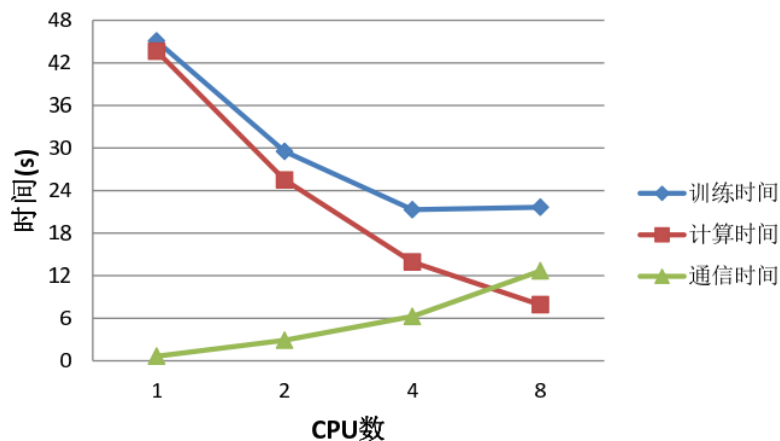
- 4 nodes, each node has dual Intel Xeon E5-2697 v3 (14 x 2 cores)
 - Infiniband QDR, Intel MKL 11.0.5.192 & MPI 4.1.1.036
- Nvidia GeForce GTX Titan
 - 2688 CUDA cores, 6GB global memory, cuBLAS

- 测试案例

- Caffe example
 - 模型: 包含6千万个参数的卷积神经网络经典模型AlexNet
 - 数据集: ImageNet数据库中, 127万张训练图片和5万张验证图片
 - **Batch-size=256** (训练过程中每次迭代处理的图片数)
 - Mini-Batch=256/n** (每个处理器完成的图片数, n是处理器数)

Evaluation

- 5120 images



n=4 s=3 **38.3%**

- communication size: 233MB/node
- Model size: 233MB

硬件	运算能力 (Tflops)	数学库	训练时间 (s)
2×E5-2697 v3	2×1.164≈2.32	MKL	29.5
4×E5-2697 v3	4×1.164≈4.65	MKL	21.3/18.9*
1×GTX770	3.2	cuBLAS	33
		cuDNN	24.3
1×K20	3.52	cuBLAS	36
1×K40	4.29	cuBLAS	26.5
		cuDNN	19.2
1×Titan	4.5	cuBLAS	26.26
		cuDNN	20.25

Summary

- CNN architecture
- Our previous work

Next Plan

- Many-core Parallelism within a single node
 - Hybrid data parallelism and model parallelism
- Heterogeneous computing
 - Intel CPU + Integrated GPU / FPGA
- How to overlap communication with computation in the distributed context
- Model resize
 - matrix decomposition or approximate matrix