Caffe

Deep learning framework by the BVLC

Created by

Yangqing Jia

Lead Developer

Evan Shelhamer

View On GitHub

Layers

To create a Caffe model you need to define the model architecture in a protocol buffer definition file (prototxt).

Caffe layers and their parameters are defined in the protocol buffer definitions for the project in caffe.proto.

Vision Layers

- Header: `./include/caffe/vision_layers.hpp`

Vision layers usually take images as input and produce other images as output. A typical "image" in the real-world may have one color channel ($c = 1$), as in a grayscale image, or three color channels ($c = 3$) as in an RGB (red, green, blue) image. But in this context, the distinguishing characteristic of an image is its spatial structure: usually an image has some non-trivial height $h > 1$ and width $w > 1$. This 2D geometry naturally lends itself to certain decisions about how to process the input. In particular, most of the vision layers work by applying a particular operation to some region of the input to produce a corresponding region of the output. In contrast, other layers (with few exceptions) ignore the spatial structure of the input, effectively treating it as "one big vector" with dimension chw.

Convolution

- Layer type: `Convolution`
- CPU implementation: `./src/caffe/layers/convolution_layer.cpp`
- CUDA GPU implementation: `./src/caffe/layers/convolution_layer.cu`
- Parameters (`ConvolutionParameter convolution_param`)
  - Required
    - `num_output` (`c_o`): the number of filters
    - `kernel_size` (or `kernel_h` and `kernel_w`): specifies height and width of each filter

  - Strongly Recommended
    - `weight_filler` [default type: `'constant'` value: `0`]

  - Optional
    - `bias_term` [default `true`]: specifies whether to learn and apply a set of additive biases to the filter outputs
    - `pad` (or `pad_h` and `pad_w`) [default 0]: specifies the number of pixels to (implicitly) add to each side of the input
    - `stride` (or `stride_h` and `stride_w`) [default 1]: specifies the intervals at which to apply the filters to the input
    - `group` (g) [default 1]: If g > 1, we restrict the connectivity of each filter to a subset of the input. Specifically, the input and output channels are separated into g groups, and the ith output group channels will be only connected to the ith input group channels.

- Input
  - `n * c_i * h_i * w_i`

- Output
  - `n * c_o * h_o * w_o`, where `h_o = (h_i + 2 * pad_h - kernel_h) / stride_h + 1` and `w_o` likewise.

- Sample (as seen in `./models/bvlc_reference_caffenet/train_val.prototxt`)

```
layer {
  name: "conv1"
  type: "Convolution"
  bottom: "data"
  top: "conv1"
  # learning rate and decay multipliers for the filters
  param { lr_mult: 1 decay_mult: 1 }
  # learning rate and decay multipliers for the biases
  param { lr_mult: 2 decay_mult: 0 }
  convolution_param {
    num_output: 96     # learn 96 filters
```

```
      kernel_size: 11    # each filter is 11x11
      stride: 4          # step 4 pixels between each filter application
      weight_filler {
        type: "gaussian" # initialize the filters from a Gaussian
        std: 0.01         # distribution with stdev 0.01 (default mean: 0)
      }
      bias_filler {
        type: "constant" # initialize the biases to zero (0)
        value: 0
      }
    }
  }
```

The `Convolution` layer convolves the input image with a set of learnable filters, each producing one feature map in the output image.

Pooling

- Layer type: `Pooling`
- CPU implementation: `./src/caffe/layers/pooling_layer.cpp`
- CUDA GPU implementation: `./src/caffe/layers/pooling_layer.cu`
- Parameters (`PoolingParameter pooling_param`)
  - Required
    - `kernel_size` (or `kernel_h` and `kernel_w`): specifies height and width of each filter

  - Optional
    - `pool` [default MAX]: the pooling method. Currently MAX, AVE, or STOCHASTIC
    - `pad` (or `pad_h` and `pad_w`) [default 0]: specifies the number of pixels to (implicitly) add to each side of the input
    - `stride` (or `stride_h` and `stride_w`) [default 1]: specifies the intervals at which to apply the filters to the input

- Input
  - $n * c * h\_i * w\_i$

- Output
  - $n * c * h\_o * w\_o$, where h_o and w_o are computed in the same way as convolution.

- Sample (as seen in `./models/bvlc_reference_caffenet/train_val.prototxt`)

```
layer {
  name: "pool1"
  type: "Pooling"
  bottom: "conv1"
```

```
    top: "pool1"
    pooling_param {
      pool: MAX
      kernel_size: 3 # pool over a 3x3 region
      stride: 2      # step two pixels (in the bottom blob) between pooling regions
    }
  }
```

Local Response Normalization (LRN)

- Layer type: `LRN`
- CPU Implementation: `./src/caffe/layers/lrn_layer.cpp`
- CUDA GPU Implementation: `./src/caffe/layers/lrn_layer.cu`
- Parameters (`LRNParameter lrn_param`)
  - Optional
    - `local_size` [default 5]: the number of channels to sum over (for cross channel LRN) or the side length of the square region to sum over (for within channel LRN)
    - `alpha` [default 1]: the scaling parameter (see below)
    - `beta` [default 5]: the exponent (see below)
    - `norm_region` [default `ACROSS_CHANNELS`]: whether to sum over adjacent channels (`ACROSS_CHANNELS`) or nearby spatial locaitons (`WITHIN_CHANNEL`)

The local response normalization layer performs a kind of "lateral inhibition" by normalizing over local input regions. In `ACROSS_CHANNELS` mode, the local regions extend across nearby channels, but have no spatial extent (i.e., they have shape `local_size x 1 x 1`). In `WITHIN_CHANNEL` mode, the local regions extend spatially, but are in separate channels (i.e., they have shape `1 x local_size x local_size`). Each input value is divided by $(1 + (\alpha/n) \sum_i x_i^2)^\beta$, where n is the size of each local region, and the sum is taken over the region centered at that value (zero padding is added where necessary).

im2col

`Im2col` is a helper for doing the image-to-column transformation that you most likely do not need to know about. This is used in Caffe's original convolution to do matrix multiplication by laying out all patches into a matrix.

Loss Layers

Loss drives learning by comparing an output to a target and assigning cost to minimize. The loss itself is computed by the forward pass and the

gradient w.r.t. to the loss is computed by the backward pass.

Softmax

- Layer type: `SoftmaxWithLoss`

The softmax loss layer computes the multinomial logistic loss of the softmax of its inputs. It's conceptually identical to a softmax layer followed by a multinomial logistic loss layer, but provides a more numerically stable gradient.

Sum-of-Squares / Euclidean

- Layer type: `EuclideanLoss`

The Euclidean loss layer computes the sum of squares of differences of its two inputs, $\frac{1}{2N} \sum_{i=1}^{N} \| x_i^1 - x_i^2 \|_2^2$.

Hinge / Margin

- Layer type: `HingeLoss`
- CPU implementation: `./src/caffe/layers/hinge_loss_layer.cpp`
- CUDA GPU implementation: none yet
- Parameters (`HingeLossParameter hinge_loss_param`)
  - Optional
    - `norm` [default L1]: the norm used. Currently L1, L2

- Inputs
  - `n * c * h * w` Predictions
  - `n * 1 * 1 * 1` Labels

- Output
  - `1 * 1 * 1 * 1` Computed Loss

- Samples

```
# L1 Norm
layer {
  name: "loss"
  type: "HingeLoss"
  bottom: "pred"
  bottom: "label"
}
```

```
# L2 Norm
layer {
  name: "loss"
  type: "HingeLoss"
  bottom: "pred"
  bottom: "label"
  top: "loss"
  hinge_loss_param {
    norm: L2
  }
}
```

The hinge loss layer computes a one-vs-all hinge or squared hinge loss.

Sigmoid Cross-Entropy

`SigmoidCrossEntropyLoss`

Infogain

`InfogainLoss`

Accuracy and Top-k

`Accuracy` scores the output as the accuracy of output with respect to target – it is not actually a loss and has no backward step.

Activation / Neuron Layers

In general, activation / Neuron layers are element-wise operators, taking one bottom blob and producing one top blob of the same size. In the layers below, we will ignore the input and out sizes as they are identical:

- Input
  - n * c * h * w

- Output
  - n * c * h * w

ReLU / Rectified-Linear and Leaky-ReLU

- Layer type: `ReLU`
- CPU implementation: `./src/caffe/layers/relu_layer.cpp`
- CUDA GPU implementation: `./src/caffe/layers/relu_layer.cu`
- Parameters (`ReLUParameter relu_param`)

   ○ Optional
      ■ `negative_slope` [default 0]: specifies whether to leak the negative part by multiplying it with the slope value rather than setting it to 0.

- Sample (as seen in `./models/bvlc_reference_caffenet/train_val.prototxt`)

```
layer {
  name: "relu1"
  type: "ReLU"
  bottom: "conv1"
  top: "conv1"
}
```

Given an input value x, The `ReLU` layer computes the output as x if x > 0 and negative_slope * x if x <= 0. When the negative slope parameter is not set, it is equivalent to the standard ReLU function of taking max(x, 0). It also supports in-place computation, meaning that the bottom and the top blob could be the same to preserve memory consumption.

Sigmoid

- Layer type: `Sigmoid`
- CPU implementation: `./src/caffe/layers/sigmoid_layer.cpp`
- CUDA GPU implementation: `./src/caffe/layers/sigmoid_layer.cu`
- Sample (as seen in `./examples/mnist/mnist_autoencoder.prototxt`)

```
layer {
  name: "encode1neuron"
  bottom: "encode1"
  top: "encode1neuron"
  type: "Sigmoid"
}
```

The `Sigmoid` layer computes the output as sigmoid(x) for each input element x.

TanH / Hyperbolic Tangent

- Layer type: `TanH`
- CPU implementation: `./src/caffe/layers/tanh_layer.cpp`
- CUDA GPU implementation: `./src/caffe/layers/tanh_layer.cu`
- Sample

```
layer {
  name: "layer"
  bottom: "in"
  top: "out"
  type: "TanH"
```

```
  }
```

The `TanH` layer computes the output as tanh(x) for each input element x.

Absolute Value

- Layer type: `AbsVal`
- CPU implementation: `./src/caffe/layers/absval_layer.cpp`
- CUDA GPU implementation: `./src/caffe/layers/absval_layer.cu`
- Sample

```
layer {
  name: "layer"
  bottom: "in"
  top: "out"
  type: "AbsVal"
}
```

The `AbsVal` layer computes the output as abs(x) for each input element x.

Power

- Layer type: `Power`
- CPU implementation: `./src/caffe/layers/power_layer.cpp`
- CUDA GPU implementation: `./src/caffe/layers/power_layer.cu`
- Parameters (`PowerParameter power_param`)
  - Optional
    - `power` [default 1]
    - `scale` [default 1]
    - `shift` [default 0]

- Sample

```
layer {
  name: "layer"
  bottom: "in"
  top: "out"
  type: "Power"
  power_param {
    power: 1
    scale: 1
    shift: 0
  }
}
```

The `Power` layer computes the output as (shift + scale * x) ^ power for each input element x.

BNLL

- Layer type: BNLL
- CPU implementation: ./src/caffe/layers/bnll_layer.cpp
- CUDA GPU implementation: ./src/caffe/layers/bnll_layer.cu
- Sample

```
layer {
  name: "layer"
  bottom: "in"
  top: "out"
  type: BNLL
}
```

The BNLL (binomial normal log likelihood) layer computes the output as log(1 + exp(x)) for each input element x.

Data Layers

Data enters Caffe through data layers: they lie at the bottom of nets. Data can come from efficient databases (LevelDB or LMDB), directly from memory, or, when efficiency is not critical, from files on disk in HDF5 or common image formats.

Common input preprocessing (mean subtraction, scaling, random cropping, and mirroring) is available by specifying TransformationParameters.

Database

- Layer type: Data
- Parameters
  - Required
    - source: the name of the directory containing the database
    - batch_size: the number of inputs to process at one time
  - Optional
    - rand_skip: skip up to this number of inputs at the beginning; useful for asynchronous sgd
    - backend [default LEVELDB]: choose whether to use a LEVELDB or LMDB

In-Memory

- Layer type: MemoryData

- Parameters
  - Required
    - `batch_size`, `channels`, `height`, `width`: specify the size of input chunks to read from memory

The memory data layer reads data directly from memory, without copying it. In order to use it, one must call `MemoryDataLayer::Reset` (from C++) or `Net.set_input_arrays` (from Python) in order to specify a source of contiguous data (as 4D row major array), which is read one batch-sized chunk at a time.

HDF5 Input

- Layer type: `HDF5Data`
- Parameters
  - Required
    - `source`: the name of the file to read from
    - `batch_size`

HDF5 Output

- Layer type: `HDF5Output`
- Parameters
  - Required
    - `file_name`: name of file to write to

The HDF5 output layer performs the opposite function of the other layers in this section: it writes its input blobs to disk.

Images

- Layer type: `ImageData`
- Parameters
  - Required
    - `source`: name of a text file, with each line giving an image filename and label
    - `batch_size`: number of images to batch together

  - Optional
    - `rand_skip`
    - `shuffle` [default false]
    - `new_height`, `new_width`: if provided, resize all images to this size

Windows

`WindowData`

Dummy

`DummyData` is for development and debugging. See `DummyDataParameter`.

Common Layers

Inner Product

- Layer type: `InnerProduct`
- CPU implementation: `./src/caffe/layers/inner_product_layer.cpp`
- CUDA GPU implementation: `./src/caffe/layers/inner_product_layer.cu`
- Parameters (`InnerProductParameter inner_product_param`)
  - Required
    - `num_output` (`c_o`): the number of filters

  - Strongly recommended
    - `weight_filler` [default `type: 'constant' value: 0`]

  - Optional
    - `bias_filler` [default `type: 'constant' value: 0`]
    - `bias_term` [default `true`]: specifies whether to learn and apply a set of additive biases to the filter outputs

- Input
  - `n * c_i * h_i * w_i`

- Output
  - `n * c_o * 1 * 1`

- Sample

```
layer {
  name: "fc8"
  type: "InnerProduct"
  # learning rate and decay multipliers for the weights
  param { lr_mult: 1 decay_mult: 1 }
  # learning rate and decay multipliers for the biases
  param { lr_mult: 2 decay_mult: 0 }
  inner_product_param {
    num_output: 1000
    weight_filler {
```

```
      type: "gaussian"
      std: 0.01
    }
    bias_filler {
      type: "constant"
      value: 0
    }
  }
  bottom: "fc7"
  top: "fc8"
}
```

The `InnerProduct` layer (also usually referred to as the <mark>fully connected layer</mark>) treats the input as a simple vector and produces an output in the form of a single vector (with the blob's height and width set to 1).

Splitting

The `Split` layer is a utility layer that splits an input blob to multiple output blobs. This is used when a blob is fed into multiple output layers.

Flattening

The `Flatten` layer is a utility layer that flattens an input of shape $n * c * h * w$ to a simple vector output of shape $n * (c*h*w)$

Reshape

- Layer type: `Reshape`
- Implementation: `./src/caffe/layers/reshape_layer.cpp`
- Parameters (`ReshapeParameter reshape_param`)
  - Optional: (also see detailed description below)
    - `shape`

- Input
  - a single blob with arbitrary dimensions

- Output
  - the same blob, with modified dimensions, as specified by `reshape_param`

- Sample

```
layer {
  name: "reshape"
  type: "Reshape"
  bottom: "input"
  top: "output"
  reshape_param {
```

```
      shape {
        dim: 0  # copy the dimension from below
        dim: 2
        dim: 3
        dim: -1 # infer it from the other dimensions
      }
    }
  }
```

The `Reshape` layer can be used to change the dimensions of its input, without changing its data. Just like the `Flatten` layer, only the dimensions are changed; no data is copied in the process.

Output dimensions are specified by the `ReshapeParam` proto. Positive numbers are used directly, setting the corresponding dimension of the output blob. In addition, ==two special values are accepted for any of the target dimension values:==

- **0** means "copy the respective dimension of the bottom layer". That is, if the bottom has 2 as its 1st dimension, the top will have 2 as its 1st dimension as well, given `dim: 0` as the 1st target dimension.
- **-1** stands for "infer this from the other dimensions". This behavior is similar to that of -1 in numpy's or `[]` for MATLAB's reshape: this dimension is calculated to keep the overall element count the same as in the bottom layer. At most one -1 can be used in a reshape operation.

As another example, specifying `reshape_param { shape { dim: 0 dim: -1 } }` makes the layer behave in exactly the same way as the `Flatten` layer.

Concatenation

- Layer type: `Concat`
- CPU implementation: `./src/caffe/layers/concat_layer.cpp`
- CUDA GPU implementation: `./src/caffe/layers/concat_layer.cu`
- Parameters (`ConcatParameter concat_param`)
  - Optional
    - `axis` [default 1]: 0 for concatenation along num and 1 for channels.

- Input
  - $n\_i * c\_i * h * w$ for each input blob i from 1 to K.

- Output
  - if `axis = 0`: $(n\_1 + n\_2 + ... + n\_K) * c\_1 * h * w$, and all input c_i should be the same.

- if `axis = 1`: $n\_1 * (c\_1 + c\_2 + ... + c\_K) * h * w$, and all input `n_i` should be the same.

- Sample

```
layer {
  name: "concat"
  bottom: "in1"
  bottom: "in2"
  top: "out"
  type: "Concat"
  concat_param {
    axis: 1
  }
}
```

The `Concat` layer is a utility layer that concatenates its multiple input blobs to one single output blob.

Slicing

The `Slice` layer is a utility layer that slices an input layer to multiple output layers along a given dimension (currently num or channel only) with given slice indices.

- Sample

```
layer {
  name: "slicer_label"
  type: "Slice"
  bottom: "label"
  ## Example of label with a shape N x 3 x 1 x 1
  top: "label1"
  top: "label2"
  top: "label3"
  slice_param {
    axis: 1
    slice_point: 1
    slice_point: 2
  }
}
```

`axis` indicates the target axis; `slice_point` indicates indexes in the selected dimension (the number of indices must be equal to the number of top blobs minus one).

Elementwise Operations

Eltwise

Argmax

`ArgMax`

Softmax

`Softmax`

Mean-Variance Normalization

`MVN`